

vexorin customer acquisition strategy

VEXORIN.AI — CUSTOMER ACQUISITION STRATEGY

Three-Tiered GTM with Target Lists, Outreach Motion, and Conversion Metrics · v1.0 · May 2026

Purpose. This document expands the three-tiered GTM frame from deck slide 9 into an operational customer acquisition strategy. Each founder-owned distribution path is documented with target customer profile, outreach motion, conversion metrics per stage, and execution responsibilities.

Audience. Tier-1 venture partners and analysts validating the customer acquisition path, plus the team's CFO and board for ongoing GTM review. The document is calibrated to be re-read quarterly with conversion-metric updates.

Strategic frame. The architecture follows from who Vexorin's founders are. So does the distribution. Each path exploits a credential specific to the founder owning it. The three paths cannot be replicated by a generalist team without recruiting three specialized founders. They compound: three independent paths to the same buyer.

1. EXECUTIVE SUMMARY

Vexorin's customer acquisition strategy operates across three distribution tiers, each owned by a specific founder and each executing a distinct motion against the same buyer set:

Tier 1 — Direct fund relationships (Amir Zada). Credibility-driven institutional sales. Converts warm relationships into LOIs into deployed Scout+Carlton instances.

Tier 2 — Language model term optimization (Peter, CTO). Technical positioning in AI-search and LLM term spaces. Surfaces Vexorin to institutional buyers conducting AI-assisted product research.

Tier 3 — Organic distribution through networked conversation platforms (Alistair Ewart). Twenty years of understanding virality algorithms applied to surfacing Vexorin within Reddit, Twitter/X, and financial-media-adjacent communities.

Together the three tiers produce **roughly 40-60% acquisition velocity advantage** over a single-channel sales motion in the same buyer category, based on the structural argument that three independent paths to the same buyer compound trust at each stage.

The Year 3 customer trajectory targets **35-40 production customers** across the underserved-tier wedge, with the Tier-1 institutional segment opening as Years 3+ moat-defending wedge.

2. CUSTOMER SEGMENTATION

2.1 Initial GTM (Years 1-3): the underserved tier

The primary customer set for Years 1-3 is the institutional segment with the same recognition problem Tier-1 funds have but without the budget to solve it internally.

Sub-segment 1 — Mid-AUM hedge funds (\$1-20B AUM). Multi-strategy, equity long/short, event-driven, and commodities funds in this AUM range. Hundreds of firms globally. Typical surveillance budget: \$100K-\$500K annually, supporting 1-3 dedicated compliance analysts.

Sub-segment 2 — Regional asset managers and boutique firms. Mid-tier asset managers with regional focus (Asia-Pacific, Latin America, Middle East), boutique investment management firms with specialized strategies. Hundreds of firms globally.

Sub-segment 3 — Emerging-manager hedge funds. Recently launched funds in their first 1-3 years of operation, typically institutional-LP-backed. Hundreds globally; high turnover but consistent overall population.

Sub-segment 4 — Family offices with institutional ambition. Single-family and multi-family offices managing \$500M-\$5B with institutional-grade infrastructure aspirations. Hundreds globally; growing segment.

Sub-segment 5 — Second-tier sovereign-wealth allocators. Sovereign-wealth allocator subsidiaries, regional pension fund operators, mid-tier government investment vehicles. Several dozen globally.

Aggregate addressable customer set across the underserved tier: several thousand institutions globally.

Year 3 capture target: 35-40 production customers (~1-2% capture rate of the addressable set).

2.2 Long-term wedge (Years 3+): Tier-1 institutions

~280 institutions globally with AUM > \$50B and active surveillance budgets > \$5M.

- **Largest hedge funds** (Citadel, Millennium, P72, Two Sigma, Renaissance, AQR, Bridgewater, DE Shaw, Citadel Wellington — top 30-50 funds globally)
- **Sovereign wealth funds** (GIC, ADIA, Norges Bank, KIA, QIA, Mubadala, Temasek, China Investment Corporation — top 50 globally)
- **G-SIBs** (JPM, GS, MS, BAML, Citi, Wells Fargo, BNY Mellon, State Street, Deutsche Bank, BNP Paribas, Crédit Agricole, Société Générale, UBS, Credit Suisse, Barclays, HSBC, Standard Chartered, Mizuho, Sumitomo Mitsui, Mitsubishi UFJ — 30 globally)

These are the moat-defending customers, not the first-revenue customers. The bridge from initial GTM to Tier-1 is the deployed customer cohort and revenue traction across three streams accumulating in front of the market, plus the validated marketplace becoming a credentialing layer Tier-1 institutions cannot ignore.

2.3 Customer profile criteria

Across both segments, the qualifying customer profile is:

- AUM or assets-under-supervision > \$1B
- Active surveillance or compliance team (even if 1-3 analysts)
- Regulatory exposure (SEC, FINRA, FCA, MAS, or equivalent)
- Active alt-data or surveillance-tooling evaluation in the past 18 months
- Procurement authority for software contracts in the \$50K-\$500K annual range (Year 1-3) or \$500K-\$5M annual range (Year 3+)

3. TIER 1 — DIRECT FUND RELATIONSHIPS (AMIR ZADA)

3.1 The motion

Amir's institutional-finance pedigree opens fund procurement conversations through warm relationships built across two decades. The motion is direct enterprise sales: outreach to a known contact at a target fund, qualifying conversation, technical evaluation introduction, LOI, deployment, expansion.

Key credentials anchoring the motion: - Director and Co-Head, New York Office, G3 Good Governance Group (current role) — institutional intelligence and cyber defense customer base - Managing Director at Renaissance Capital and Exotix Ltd (prior) — frontier-markets investment banking - Senior Managing Director at Global Frontera Markets (prior) — frontier-markets investment banking leadership - MSc Management, London School of Economics; MSc International Management, HEC Paris - Member, Board of Advisors at TLG Capital - Active on professional networks (LinkedIn, X under @Zada7) - Recruited into British intelligence from high school (intelligence-community pedigree)

3.2 Target customer list (Year 1)

Direct outreach targets — institutions where Amir has a known contact or established relationship:

- **Mid-tier hedge funds** in NYC, London, Hong Kong, Singapore where Amir's banking history connects to current fund principals
- **Frontier-market focused funds** (the Renaissance / Exotix / Global Frontera customer base) that maintain operations in markets where institutional manipulation is most prevalent
- **TLG Capital portfolio companies and adjacent firms** — TLG's investor and operating network
- **G3 Good Governance Group's institutional customer base** (where appropriate — distinct from Vexorin's surveillance offering, but adjacent in intelligence-and-cyber-defense category)
- **Frontier-market sovereign-wealth secondary tiers** (Middle East, Asia-Pacific) where Amir has historical relationships

Specific target list maintained internally — approximately 30-50 named first-degree contacts across the underserved-tier wedge.

3.3 Outreach motion

Stage 1 — Warm outreach via direct network. Amir reaches out to known contacts at qualified target institutions. Initial conversation focuses on (a) the recognition problem (institutional manipulation that escapes single-surface surveillance), (b) Vexorin’s substrate, (c) scope of relevant deployment.

Stage 2 — Qualifying conversation. Once initial interest is established, qualifying conversation determines: (1) the customer’s current surveillance posture, (2) the customer’s data infrastructure (relevant to deployment), (3) the customer’s procurement authority and budget, (4) the timeline for evaluation. Approximately 30-50% of Stage 1 contacts proceed to Stage 2.

Stage 3 — Technical evaluation introduction. CTO joins the conversation for technical evaluation. Scout walkthrough (45 minutes), Carlton walkthrough, BRAIN architecture overview, data-handling architecture review. Approximately 50-70% of Stage 2 conversations proceed to Stage 3.

Stage 4 — LOI. Letter of Intent commits to deployment evaluation under specified terms (pricing model, data-handling architecture, deployment timeline). Approximately 50-70% of Stage 3 evaluations produce LOIs. Currently several LOIs in pipeline.

Stage 5 — Deployment. Scout and Carlton deploy on-prem in 2-3 days. Customer begins paying \$50K platform fee (Stream 1), agent consumption (Stream 2). CALIBRATE begins customer-specific FPR recalibration. Approximately 70-90% of LOIs convert to deployment.

Stage 6 — Expansion. Customer activates additional agents as they come online (ECHO, ATLAS, CIPHER from Month 5; remaining agents from Month 8+). Customer integrates with Validated Data Marketplace post-launch.

3.4 Year 1 conversion targets

Stage	Activity	Target volume	Conversion to next stage
1	Warm outreach	50 contacts	40-50%
2	Qualifying conversation	20-25 conversations	50-70%
3	Technical evaluation	10-17 evaluations	50-70%
4	LOI	5-12 LOIs	70-90%
5	Deployment	3-10 deployments (Year 1 target: 5-7)	n/a

3.5 Year 3 trajectory (Tier 1 contribution)

- **Cumulative deployments via Tier 1 by Year 3:** 15-20 customers
- **Average revenue per Tier-1-acquired customer at Year 3:** ~\$500K (platform fee + agent consumption + early marketplace participation)
- **Cumulative Tier 1 revenue contribution at Year 3:** \$7.5M-10M ARR (out of \$5.3M base case across all three tiers — Tier 1 is the largest revenue contributor, with Tiers 2 and 3 amplifying Tier 1 conversion velocity rather than producing distinct customer cohorts)

4. TIER 2 — LANGUAGE MODEL TERM OPTIMIZATION (PETER, CTO)

4.1 The motion

The 2026 institutional buyer increasingly discovers products through AI-assisted search rather than traditional Google search. They query Claude, ChatGPT, Perplexity, or institutional-tier AI search products with questions like:

- “What surveillance products integrate with our existing infrastructure?”
- “Compare manipulation detection vendors for a mid-tier hedge fund”
- “AI-based market manipulation detection on-prem deployment options”
- “Multi-surface surveillance vs single-surface surveillance — which to choose?”

Vexorin’s appearance in those answers depends on three factors: (1) the structured documentation visible in the technical commons (GitHub, arXiv, SSRN, technical blog posts, white papers), (2) the LLM-aware product positioning (claims structured to be retrievable in agentic search), (3) the presence of Vexorin in adjacent technical conversations (Stack Overflow-class technical Q&A, GitHub technical discussions, arXiv paper references).

Peter (CTO) owns this motion because the technical authenticity required is engineering-deep. The motion does not work if the surfaced content is marketing-grade; it works only when the technical documentation is rigorous enough that an AI-assisted research tool would surface it as authoritative.

4.2 The output channels

Channel 1 — Vexorin technical commons. Public technical documentation hosted on Vexorin's domain: BRAIN architecture overview, agent technical specifications (with appropriate redactions for trade-secret protection), citation-audit summary, methodology blog posts, integration guides for institutional buyers.

Channel 2 — Open-source contributions. Engineering team contributions to relevant open-source projects in the agentic-infrastructure category (MCP protocol implementations, signal-aggregation primitives, audit-trail infrastructure for AI systems). Each contribution is a citation-worthy reference point that surfaces in AI-assisted research.

Channel 3 — Technical blog content. Engineering blog posts on relevant topics: multi-surface vs single-surface surveillance architecture, conviction-aggregation methodology, falsifiability in AI-driven detection, on-prem AI deployment patterns, customer-specific calibration in financial-services AI. Each post is structured to be retrievable in AI-assisted search.

Channel 4 — Research paper references. Vexorin publishes a small number of high-quality research notes on the methodology (cross-validation results from Q3 2026, family taxonomy validation, geometric manipulation recognition methodology). The notes are not academic papers — they are research-grade summaries that anchor the architecture's intellectual depth. Hosted on Vexorin's domain with citation metadata that AI search tools can surface.

Channel 5 — Adjacent technical Q&A. Engineering team participation in technical Q&A communities relevant to the institutional-AI-deployment category. Stack Overflow-class technical answers, GitHub technical discussions, arXiv paper review activity. The participation is authentically technical, not marketing-driven.

4.3 Year 1 execution

Q1 (Months 1-3 post-close): - Vexorin technical commons live (Channel 1): BRAIN architecture overview, agent technical specifications, citation-audit summary, integration guide - First five engineering blog posts (Channel 3) - First open-source contribution (Channel 2)

Q2 (Months 4-6): - First research note published (Channel 4): cross-validation results from the Q1 analytical pass - Continued technical commons expansion - Five additional blog posts

Q3 (Months 7-9): - Second research note: family taxonomy validation - Quarterly review of AI-search retrievability (test queries against Claude, ChatGPT, Perplexity to verify Vexorin appears in relevant answers)

Q4 (Months 10-12): - Third research note: methodology summary (geometric manipulation recognition) - AI-search retrievability hits target (Vexorin appears in 50%+ of relevant institutional buyer queries on AI-assisted search)

4.4 Conversion metrics

The Tier-2 conversion metric is **AI-search retrievability**, measured quarterly through a panel of institutional-buyer-style queries.

- **Quarterly query panel:** ~20 representative queries an institutional buyer might ask Claude / ChatGPT / Perplexity when researching surveillance products
- **Target retrievability:** Vexorin appears in answer for 50%+ of relevant queries by Q4
- **Surfacing format:** authoritative reference (cited as primary source for at least 30% of relevant answers, vs. tertiary mention)

Tier 2 does not produce direct deployments. It amplifies Tier 1 by ensuring that when Amir's outreach reaches a target fund's research analyst, the analyst's AI-assisted research already surfaces Vexorin as authoritative. The conversion velocity advantage from Stage 2 to Stage 3 (qualifying conversation to technical evaluation) is approximately 2x for buyers who arrive at Vexorin already familiar with the architecture.

5. TIER 3 — ORGANIC DISTRIBUTION (ALISTAIR EWART)

5.1 The motion

Twenty years of understanding virality algorithms and how to optimize them translate directly into Vexorin's organic distribution. Surfacing Vexorin within the networked conversations on Reddit, Twitter/X, and the financial-media-adjacent communities where institutional buyers are increasingly informed.

The principle: institutional buyers no longer rely solely on direct vendor outreach or trade publications. They participate in networked conversations — finance Twitter, specialized Reddit communities (r/SecurityAnalysis, r/quantfinance, niche compliance subreddits), Substack newsletters, podcast ecosystems — that shape their awareness of new infrastructure categories. A founder who understands how attention compounds across these networks can surface Vexorin within them organically, building category awareness that converts into procurement-stage interest.

Alistair owns this motion because the operational expertise required is exactly what KOAT.AI was built on: understanding how networked conversations propagate, where attention concentrates, how communities respond to new content, and how to position content authentically rather than as marketing-spam.

5.2 The output channels

Channel 1 — Founder-led commentary on networked conversation platforms. Alistair's authentic participation in finance Twitter / X (under @AlistairEwart or equivalent), Reddit (relevant subreddits), and specialized finance communities. Commentary on manipulation events as they emerge, on category trends, on relevant academic papers. Authentic founder voice; not marketing.

Channel 2 — Substack / podcast ecosystem participation. Alistair as guest on relevant finance and AI-infrastructure podcasts (target: 10-15 podcast appearances Year 1). Substack newsletter focused on manipulation recognition, market microstructure, and AI in financial services (subscriber acquisition target: 5,000+ by Year 1).

Channel 3 — Conference and event participation. Speaking engagements at relevant industry conferences (FinTech Connect, AI in Finance Summit, RegTech & Compliance NYC, DC Quant Conference, equivalent international events). Target: 4-6 speaking engagements Year 1.

Channel 4 — Vexorin's own owned media. Vexorin Substack, Vexorin podcast (interview format with industry guests), Vexorin YouTube content (technical explainers, architecture walkthroughs). Each owned-media artifact is engineered for organic distribution within the networks Alistair operates in.

5.3 Strategic discipline

The motion only works if the participation is authentic. Three operational disciplines:

Discipline 1 — No spam. Vexorin does not promote itself unsolicited in communities. Alistair's participation is genuine (commentary on emerging events, contributions to discussions); Vexorin's product surface emerges organically when relevant.

Discipline 2 — High signal-to-noise ratio. Every piece of content (commentary, post, podcast appearance, conference talk) has substantive signal. Marketing-grade content destroys the credibility the motion depends on.

Discipline 3 — Founder authenticity. Alistair owns the voice. The motion does not work if Alistair's voice is replaced by a marketing or PR team. Authenticity is the moat; if it dilutes, the motion dilutes with it.

5.4 Year 1 conversion targets

The Tier-3 conversion metric is **inbound interest velocity** — institutional buyers who reach out to Vexorin (or to Amir directly) without prior direct outreach, citing organic discovery (heard on a podcast, read on Substack, saw a Twitter conversation, etc.).

- **Quarterly inbound interest target:** 10-15 inbound qualified inquiries by Q4
- **Cumulative inbound by Year 1:** 30-50 qualified inbound inquiries
- **Conversion to Stage 2 (qualifying conversation):** 30-50% of inbound (10-25 qualifying conversations)

Tier 3 amplifies Tier 1 in the same way Tier 2 does — buyers who arrive familiar with Vexorin from organic distribution convert through subsequent stages 1.5-2x faster than buyers arriving cold.

6. TIER COMPOUNDING

The three tiers are not parallel; they compound.

The three-step compound: Amir's direct outreach reaches a target fund's research analyst. Peter's LLM term optimization makes that analyst find Vexorin when they query Claude or ChatGPT in their research workflow. Alistair's organic distribution makes the same analyst already familiar with Vexorin from networked conversations before the procurement conversation begins.

The compound effect on Stage 3 → Stage 4 conversion:

- Without Tiers 2/3: 50-70% conversion (institutional baseline)
- With Tiers 2 active (analyst's AI-assisted research surfaces Vexorin): 65-85% conversion
- With all three tiers active (analyst arrives already familiar with Vexorin): 75-90% conversion

The compound is the single largest GTM advantage Vexorin has over single-channel competitors. Each individual tier produces meaningful direct customer acquisition; the compound produces an acquisition velocity that justifies the founder bandwidth investment in three concurrent channels.

7. EXECUTION RESPONSIBILITIES AND CADENCE

7.1 Ownership

Tier	Owner	Time allocation
Tier 1 — Direct fund relationships	Amir Zada	60-70% of his Vexorin time
Tier 2 — LLM term optimization	Peter (CTO)	15-20% of his Vexorin time (technical commons production)
Tier 3 — Organic distribution	Alistair Ewart	30-40% of his Vexorin time

7.2 Review cadence

Weekly: Founder GTM standup (each tier’s pipeline review, immediate follow-up commitments). **Monthly:** GTM review with CFO (conversion metrics per stage, revenue trajectory tracking, budget against use-of-funds GTM allocation). **Quarterly:** Board-level GTM review (tier-level performance, strategic adjustments, pipeline health).

7.3 Crisis cadence

If conversion at any stage falls more than 30% below target for two consecutive months:

1. Founder GTM emergency review (full-day session with all three founders, CFO, CTO).
2. Diagnosis of bottleneck (target market, messaging, pricing, technical fit, or competitive landscape).
3. Adjustment to tier-level execution. Potential mid-quarter board update.

8. EXPANSION INTO TIER-1 INSTITUTIONS (YEARS 3+)

The expansion path from underserved-tier to Tier-1 has four sequential gates:

Gate 1 — Production deployment cohort. 12+ underserved-tier customers in production with documented revenue traction across all three streams. Target: Year 2 end.

Gate 2 — Validated marketplace credentialing. 25+ data providers onboarded with active revenue share flowing. Marketplace becomes credentialing layer Tier-1 institutions cannot ignore. Target: Year 2-3.

Gate 3 — Third-party validation. Cross-validation results published, walk-forward log demonstrating real-time detections beyond Iran 2026, regulatory engagement at SEC FinHub or FINRA OFI level. Target: Year 2-3.

Gate 4 — Tier-1 procurement engagement. Formal procurement conversations at Tier-1 hedge funds, sovereign wealth funds, G-SIBs. Initial contracts likely 18-24 months from first conversation given Tier-1 procurement timelines. Target: Year 3-5.

The Tier-1 wedge converts at materially higher ACV — typical Tier-1 production contracts in the \$1-5M annual range, vs. \$250-500K for underserved-tier production. Tier-1 customer cohort drives the Series B and Series C revenue thesis.

9. VERSION CONTROL

Version	Date	Changes
1.0	2026-05-07	Initial Customer Acquisition Strategy. Three-tiered GTM with founder ownership, target customer segmentation, outreach motion, conversion metrics, tier compounding, execution cadence, Tier-1 expansion path.