

# vexorin data strategy compliance position

---

## VEXORIN.AI — DATA STRATEGY & COMPLIANCE POSITION

---

### Strategic Analysis for the Data Room · v1.0 · May 2026

---

**Purpose.** This document presents Vexorin's data strategy and regulatory compliance position, calibrated to the regulatory regimes the company will operate under and anchored to comparable-company precedent that establishes tractability of the compliance challenges. **This is strategic analysis, not legal opinion.** Counsel review and validate before any customer contracts are signed; legal opinion on specific regulatory questions is engaged separately.

**Audience.** Senior partners and analysts at Tier-1 venture firms reviewing Vexorin's data room. The document's job is to demonstrate that (1) Vexorin's data and compliance posture is credibly grounded in the same architecture comparable companies operate under, (2) the regulatory regimes Vexorin will face have established compliance precedents in adjacent categories, and (3) Vexorin's specific architectural choices pre-empt several categories of exposure that less-thoughtful competitors face.

**Companion documents.** This document complements the BRAIN Litepaper (technical architecture), Master Reference (commercial and strategic positioning), and the Competitive Landscape Analysis (competitor-by-competitor depth). Where this document references regulatory regimes or comparable companies, primary sources are footnoted at the end.

---

## 1. EXECUTIVE SUMMARY

---

Vexorin's data strategy follows three architectural choices that materially constrain regulatory exposure relative to comparable AI-deployed financial-technology platforms:

**Choice 1 — On-prem deployment as the default.** Customer data does not leave the customer's environment. This architectural choice eliminates entire categories of cross-border data-flow exposure that hosted-SaaS platforms face. Palantir Foundry has built an \$80B+ market capitalization company on the same architectural choice; the deployment model is industry-validated.

**Choice 2 — Public-data-only ingestion at the architecture level.** Scout ingests publicly-available social media posts, news articles, and public web content. The ingestion model parallels Dataminr's licensing posture, which has scaled to global Tier-1 financial institution customers under the same regulatory regimes Vexorin will face.

**Choice 3 — Manipulation detection, not credit scoring or biometric identification.** Vexorin's specific use case sits in a regulatory category the EU AI Act treats favorably. **Fraud detection is explicitly excluded from the high-risk credit-scoring classification under Annex III** (Annex III points 5b/5c). This positioning materially reduces Vexorin's regulatory exposure relative to AI fintechs operating in credit, insurance, or biometric verification categories.

The three choices compose to a single thesis: **Vexorin's compliance challenges are tractable because comparable companies have established the precedent for tractability.** This document provides the sourced foundation for that thesis.

---

## 2. THE DATA INGESTION ARCHITECTURE

---

### 2.1 What Vexorin ingests

The Vexorin platform ingests data across multiple classes corresponding to its agent set. For Deck 1's data room scope, the relevant data classes are those associated with **production-ready agents** (Scout, Carlton) plus the architectural commitments for the additional agents the round funds.

**Scout — public conversation data.** Real-time monitoring across major social platforms, 20,000+ news sites, 100+ languages. Coverage is restricted to **publicly-available content** by architectural design: posts, articles, and public web content visible to any user without authentication. No private content, no scraped DM content, no behind-paywall content unless explicitly licensed.

**Carlton — public market data.** Stock-price time series, volume, options chain data sourced from licensed market-data

feeds (NYSE, NASDAQ, OPRA composite, CME Group, ICE). Standard institutional market-data licensing posture; no novel exposure beyond what every quantitative trading platform faces.

**Future agents (post-pre-seed engineering build).** ECHO (options/vol/skew data, same market-data class as Carlton). ATLAS (regulatory filings — EDGAR, SEDAR, FCA — all public-record sources). CIPHER (insider-trading filings — Form 4, STOCK Act disclosures — all public-record sources). NEXUS (cross-asset market data). HERALD (event-study public corporate disclosures). ARGUS (geopolitical news and public policy data). MAGNUS (public central-bank communications). ORION (public FX market data). AXIOM (regime classification from observed market structure — derived from public market data).

**Critical architectural property: Vexorin's data ingestion is structurally restricted to public-record and licensed-market-data sources.** No private personal data. No biometric data. No private financial records. No protected categories under GDPR Article 9.

## 2.2 Customer data — handled separately, never aggregated

Customer data (a fund's order flow, portfolio holdings, internal compliance communications) is processed **on-prem within the customer's own environment**. Vexorin's software runs in the customer's infrastructure; customer data does not flow back to Vexorin's central infrastructure. This is the same architectural model Palantir Foundry deploys for its government and financial-services customers, with the same data-sovereignty implications.<sup>1</sup>

The architectural separation is operationally important: customer data and Vexorin's public-data ingestion are two distinct data planes that meet only at the analytical layer (where BRAIN aggregates signals from public-data agents with customer-data signals to produce conviction scores). Cross-flow between the planes happens only inside the customer's environment, under the customer's data governance regime.

## 2.3 The on-prem default

Vexorin's BRAIN Litepaper Section 13 documents three deployment models: standard hosted, self-hosted (default for compliance-sensitive customers), and air-gapped (default for regulator and major-bank deployments). For Vexorin's primary customer demographic — institutional financial customers — the **on-prem self-hosted model is the default deployment posture**. Hosted deployment is offered for customers without specific data-residency or sovereignty constraints (smaller funds, proof-of-concept evaluations).

The on-prem default is a deliberate strategic and regulatory choice. It is the same architectural choice that allowed Palantir to scale into government and Tier-1 financial customers despite the regulatory complexity of those segments. Palantir Foundry deploys via on-premise/hybrid configurations and supports HIPAA, GDPR, and FedRAMP High compliance through this architecture.<sup>2</sup> Vexorin inherits the same architectural advantages from the same architectural choice.

---

# 3. THE REGULATORY FRAMEWORK

---

## 3.1 EU AI Act — the load-bearing 2026 regime

The EU AI Act (Regulation (EU) 2024/1689) entered into force in August 2024, with phased compliance milestones running through 2027. The **August 2, 2026 deadline** is when high-risk AI system requirements become fully enforceable for financial services.<sup>34</sup>

For AI systems classified as "high-risk" under Annex III, the Act imposes substantial obligations including continuous risk management (Article 9), data governance and bias monitoring (Article 10), technical documentation (Article 11), audit-trail logging (Article 12), human oversight (Article 14), and conformity assessment (Article 43).<sup>5</sup> **Penalties for non-compliance reach €15 million or 3% of global annual turnover for general infringements, and €35 million or 7% of turnover for prohibited practices.**<sup>6</sup>

**Vexorin's strategic positioning under the AI Act.** The Annex III high-risk categories most relevant to financial services are credit scoring (point 5b), insurance pricing (point 5c), and biometric identification (point 1). Vexorin operates in **none of these categories**. Vexorin's use case is **manipulation detection** — closer in regulatory positioning to fraud detection than to credit scoring.

**This positioning matters.** The matproof.com analysis of EU AI Act fintech compliance explicitly notes: "*Fraud detection AI is explicitly excluded from the high-risk credit scoring category.*"<sup>7</sup> The relevant Annex III text classifies as high-risk those AI systems used for "*creditworthiness assessment and credit scoring of natural persons*" — assessing an individual's likelihood of repaying credit. Vexorin assesses whether observed market-firing patterns match documented manipulation geometries; it does not assess any natural person's creditworthiness.

This does not mean Vexorin is unregulated under the AI Act — the Act's general-purpose AI obligations (transparency, technical documentation, copyright compliance, applied from August 2025<sup>8</sup>) still apply, and Vexorin's customers may be

subject to high-risk obligations in their own AI deployments. Vexorin's posture is to comply with general-purpose AI obligations directly and to architect the platform so that customers can satisfy their own high-risk obligations using Vexorin's outputs.

**Architectural advantages relative to high-risk obligations.** Even though Vexorin is not classified as high-risk, the BRAIN architecture is deliberately built to support customers who are operating high-risk systems and who need their AI infrastructure to support their compliance. Specifically:

- **Article 12 audit-trail logging.** BRAIN's three-layer auditable provenance (L1 summary, L2 agent breakdown, L3 forensic trace) directly satisfies the audit-trail requirements that high-risk-system deployers must meet.<sup>9</sup>
- **Article 14 human oversight.** BRAIN is "not a black box (intentionally)" per BRAIN Litepaper Section 14.1 — outputs are explanations, not autonomous decisions; customer analysts review every high-conviction case before regulatory reporting or other action.
- **Article 13 explainability.** BRAIN's outputs trace back to source signals; every signal traces to weighted academic priors; every prior is documented and replicable. The system is built to satisfy the explainability requirement that the EU AI Act imposes on high-risk systems.

**The strategic implication.** Customers deploying Vexorin to support their own high-risk AI obligations face a more credible compliance path than they would with black-box ML competitors. A bank that is itself classified as a high-risk-system deployer can use Vexorin's audit trail to demonstrate compliance with Articles 12-14 in a way they cannot with opaque competitor systems.

### 3.2 GDPR — the long-standing baseline

The General Data Protection Regulation (Regulation (EU) 2016/679) has been in force since 2018 and applies to any processing of personal data of individuals in the EU. Vexorin's exposure under GDPR is materially constrained by architectural choices.

**Vexorin processes minimal personal data by architectural design.** Scout's ingestion is restricted to publicly-available content. Public social media posts contain personal data (the author is identifiable in many cases), but processing of public content is broadly permitted under GDPR's "legitimate interest" basis (Article 6(1)(f)) when (i) the processing is necessary, (ii) the legitimate interest is not overridden by the data subject's rights, and (iii) the data subject has made the content public. Dataminr operates under this same architecture and has scaled to global Tier-1 financial institutional customers under GDPR.<sup>1011</sup>

**Customer data handling is governed by the customer.** Customer order flow, portfolio data, and internal communications are processed on-prem under the customer's own GDPR controls. Vexorin acts as a data processor (or sub-processor) under Article 28; the relationship is governed by Data Processing Agreements with audit rights, sub-processor approval, and incident reporting clauses standard in fintech vendor agreements.

**Right-to-be-forgotten architecture.** Scout's ingestion respects platform-level deletion: when a social media user deletes a public post, the post is removed from Scout's monitoring scope on the next ingestion cycle. This mirrors Dataminr's posture, which states explicitly: *"Social media users who do not want their data processed by Dataminr can make their account or posts private."*<sup>12</sup>

### 3.3 DORA — the financial-services digital resilience regime

The Digital Operational Resilience Act (Regulation (EU) 2022/2554) became fully applicable on January 17, 2025 and imposes ICT risk management, third-party risk management, and incident-reporting obligations on financial entities operating in the EU.<sup>13</sup>

**Vexorin's exposure under DORA is as a third-party ICT service provider** to financial-entity customers. DORA's third-party ICT risk management framework (Article 28-30) requires customers to manage their critical ICT third-party providers under contractual frameworks that include audit rights, incident reporting, sub-contracting controls, and exit strategies. Vexorin's customer contracts will include these clauses standard.

The Digital Omnibus proposal introduced by the European Commission in November 2025 simplifies and harmonizes the AI Act, GDPR, NIS2, DORA, and the Data Act with a single incident-reporting point and aligned breach-notification thresholds.<sup>14</sup> Vexorin will track Digital Omnibus implementation and integrate the harmonized incident-reporting framework into its compliance posture as the regime stabilizes.

### 3.4 US regulatory regimes

**SEC and FINRA.** Vexorin is not a registered broker-dealer, investment adviser, or exchange. Vexorin's outputs are evidence for customer analyst decisions; the customer is the regulated entity that uses those outputs in their own surveillance and compliance workflows. Vexorin's architectural decision to position itself as a detection tool for institutional buyers (per BRAIN Litepaper Section 12.5) — not as a regulator-targeted product — means Vexorin's direct

regulatory engagement with SEC FinHub or FINRA OFI is not a near-term GTM priority. The architecture is built to **support** customer regulatory workflows; Vexorin is not itself a regulated entity.

**State privacy laws.** California's CCPA/CPRA, Virginia's VCDPA, Colorado's CPA, Connecticut's CTDPA, and the growing patchwork of state privacy laws each impose obligations on businesses processing personal data of state residents. Vexorin's exposure is constrained by the same architectural choices that limit GDPR exposure — public-data-only ingestion, on-prem customer-data handling, minimal personal data processing.

**US Treasury / OFAC.** Vexorin's customer-screening posture follows standard institutional-vendor practice: customer onboarding includes OFAC screening; sanctioned entities are not onboarded; the customer base is geographically vetted at contract initiation.

### 3.5 UK regulatory regime (post-Brexit)

The UK's AI regulatory framework is developing under the principles-based approach announced in the 2023 white paper, with sector-specific regulators (FCA, PRA, ICO, CMA) implementing AI guidance within their existing remits. UK financial-services AI guidance from the FCA emphasizes accountability, explainability, and human oversight — principles BRAIN's architecture is built around.

Vexorin's UK exposure is governed by UK GDPR (mirroring EU GDPR with minor divergences), the Financial Services and Markets Act 2000, and the FCA's evolving AI guidance. The UK does not exempt Vexorin from the EU AI Act for activities serving EU customers from a UK base.<sup>15</sup>

### 3.6 Cross-border data flows

Vexorin's on-prem deployment architecture eliminates most cross-border data-flow questions because customer data does not flow across borders — it stays in the customer's environment. The exception is Vexorin's central infrastructure (which hosts the federated registry, the validated data marketplace, and operational support for customers) — this infrastructure resides in jurisdictions with adequate data-protection regimes (initial deployment in Canada and the US, with EU regional infrastructure planned for European-customer support).

For customers requiring strict data residency (EU customers, regulated US customers, Asian sovereign-wealth customers), the air-gapped deployment model offered per BRAIN Litepaper Section 10.1.3 eliminates cross-border data flow entirely.

---

## 4. PRECEDENT FROM COMPARABLE COMPANIES

### 4.1 Dataminr — the public-conversation data precedent

Dataminr is the closest precedent for Scout's data layer. Dataminr ingests publicly-available data from social media, news sites, blogs, and official reports; processes it via AI/ML to detect emerging events; and serves the output via subscription to institutional customers including Tier-1 financial institutions, government agencies, and crisis-management teams.<sup>16</sup>

**Dataminr's compliance posture, summarized from public sources:** - Public-data-only ingestion model, with explicit privacy policy stating publicly-available information is the basis of service delivery.<sup>17</sup> - Subscription-based commercial model with annual or multi-year contracts to institutional customers — same commercial model Vexorin's Stream 1 + Stream 2 follows.<sup>18</sup> - Compliance posture explicitly addresses DORA and US state privacy laws as ongoing operational concerns.<sup>19</sup> - Third-party content licensing addresses platform-specific terms (e.g., the X/Twitter Developer Terms) that constrain how platform content can be used downstream.<sup>20</sup>

**Implication for Vexorin.** Dataminr has demonstrated that the public-conversation data ingestion model scales to Tier-1 financial-institutional customers under EU and US regulatory regimes. The compliance challenges Dataminr navigates (DORA, GDPR, US state privacy laws) are the same challenges Vexorin will navigate. **The path is established.** Vexorin's commercial scaling does not depend on solving compliance problems that are categorically novel.

### 4.2 Palantir Foundry — the on-prem AI deployment precedent

Palantir Foundry is the dominant precedent for Vexorin's on-prem deployment architecture. Foundry supports flexible deployment including cloud-based SaaS and on-premise/hybrid configurations, with extensive compliance certifications including HIPAA, GDPR, and FedRAMP High.<sup>21</sup> The platform's security framework — granular access controls, mandatory encryption at rest and in transit, multi-factor authentication, comprehensive audit logging — establishes the institutional-grade baseline for AI platforms deploying into compliance-sensitive customer environments.

**Implication for Vexorin.** Palantir has scaled to a multi-tens-of-billion market capitalization on the same architectural choice (on-prem deployment with customer data sovereignty) that Vexorin's BRAIN Litepaper Section 13 commits to. The

customer-data-stays-customer-side architecture is industry-validated; the certification regimes Palantir holds (FedRAMP High, HIPAA, GDPR) provide the path Vexorin will follow as it scales into customer segments that require formal certification. Vexorin's pre-seed stage does not require these certifications immediately, but the architectural choice that enables them is in place from day one.

### 4.3 Adjacent precedents

**Snowflake — data clean-room architecture.** Snowflake's data clean rooms allow multiple parties to query shared data without revealing the underlying data to each other. The architectural pattern is relevant to Vexorin's Validated Data Marketplace (BRAIN Litepaper Section 20) where third-party data providers will validate against Vexorin's mesh without exposing their proprietary data structures.

**NICE Actimize, Behavox, SteelEye — financial communications surveillance precedents.** These competitors operate in adjacent categories (compliance communications surveillance) under the same regulatory regimes Vexorin will face. Their existence and scaling within the regulatory framework establishes that financial-services AI surveillance is operationally viable; Vexorin's architectural advantages over these competitors (multi-surface vs single-surface) are competitive, not regulatory.

**Anduril — intelligence-community-grade AI infrastructure.** Anduril's deployment architecture for defense and intelligence customers establishes the precedent for AI infrastructure deployed into customer environments with the highest data-sovereignty requirements. Given the team's intelligence-community pedigree (Amir's British intelligence background, KOAT.AI's government-deployment provenance), the Anduril precedent is particularly relevant for Vexorin's longer-term sovereign-customer deployments.

---

## 5. SPECIFIC ARCHITECTURAL CHOICES AND THEIR REGULATORY IMPLICATIONS

---

### 5.1 The “no black box” architectural choice

BRAIN does not use deep neural networks for its core conviction-aggregation logic (per BRAIN Litepaper Section 14.1). BRAIN uses ML in specific bounded contexts within agents (e.g., ATLAS's RUSBoost classifier, SCOUT's bot-detection clustering), but the aggregation logic is deliberately rule-based and traceable.

**Regulatory implication.** The EU AI Act's Article 13 explainability requirements and the FCA's UK AI guidance both emphasize the need for AI-assisted decisions to be auditable and explainable. BRAIN's architecture is **structurally aligned with this requirement** — the rule-based aggregation produces explanations natively, not as post-hoc retrofits. Customers using Vexorin do not face the explainability problem that customers using black-box ML systems face under high-risk AI deployment scenarios.

### 5.2 The “research-derived priors” disclosure choice

BRAIN's day-one weights are research-derived priors from the 50-case enforcement library, not validated posteriors (BRAIN Litepaper Section 14.6). This disclosure is honest, pre-registered, and carried throughout customer-facing materials.

**Regulatory implication.** Article 13 of the EU AI Act requires high-risk-system providers to provide users with information that allows them to interpret the system's output. By explicitly framing BRAIN's day-one performance as research-derived rather than customer-validated, Vexorin pre-empts a category of disclosure failure that affects AI vendors who claim performance levels their day-one systems cannot deliver. The honest framing is also the compliant framing.

### 5.3 The effect-size compression choice

Vexorin applies effect-size compression of 30-60% to all weights at deployment, reflecting post-publication compression in academic-finance signal literature (Prior Weight Matrix Part 6.2).

**Regulatory implication.** Article 10 of the EU AI Act requires data governance practices that ensure training, validation, and testing data are relevant, sufficiently representative, and free from errors to a degree appropriate to the intended purpose. The effect-size compression policy is the operational expression of this requirement — Vexorin's weights reflect realistic 2026-era effect sizes rather than optimistic literature magnitudes, which materially reduces the bias-and-error risk that Article 10 targets.

### 5.4 The customer-pull extension layer choice

Both extension layers (federated registry, validated data marketplace) operate on customer-pull architecture: nothing arrives in a customer environment without the customer browsing the registry from inside their own deployment and choosing to integrate (BRAIN Litepaper Section 20.4).

**Regulatory implication.** This choice eliminates several categories of contractual and regulatory exposure that vendor-push extension architectures face. Customer consent is built into the architecture; supply-chain transparency is automatic; the customer's existing third-party-risk-management framework governs every integration.

---

## 6. KEY DUE-DILIGENCE QUESTIONS, ANSWERED

---

The following questions are the most likely to surface in a Tier-1 partner's data-room review of Vexorin's compliance position. Each is answered with the relevant architectural choice and sourced precedent.

**Q1 — Is Vexorin classified as high-risk under the EU AI Act?** A — No. Vexorin's manipulation detection use case does not fall under any Annex III high-risk category. Fraud detection is explicitly excluded from the high-risk credit-scoring category (Annex III point 5b/5c).<sup>22</sup> Vexorin is subject to the Act's general-purpose AI obligations applied from August 2025, which the platform is built to satisfy through its rule-based aggregation architecture and three-layer audit trail.

**Q2 — How does Vexorin handle customer data?** A — On-prem within the customer's environment. Customer data does not flow back to Vexorin's central infrastructure. The architectural model parallels Palantir Foundry's deployment posture for government and financial-services customers, which has supported HIPAA, GDPR, and FedRAMP High compliance.<sup>23</sup>

**Q3 — How does Vexorin handle social media and news data?** A — Public-data-only ingestion, structurally restricted at the architecture level. The model parallels Dataminr's licensing posture, which has scaled to Tier-1 financial-institutional customers under EU and US regulatory regimes.<sup>2425</sup> Right-to-be-forgotten is supported through ingestion-cycle removal of deleted public content.

**Q4 — What's Vexorin's exposure to GDPR?** A — Constrained by architectural choices: public-data-only ingestion, on-prem customer-data handling, no protected-category data processing. Customer-side data processing operates under Article 28 sub-processor agreements with standard audit, sub-processor approval, and incident-reporting clauses.

**Q5 — What about DORA?** A — Vexorin's exposure is as a third-party ICT service provider to financial-entity customers. DORA's third-party risk-management framework applies; Vexorin's customer contracts will include the standard DORA clauses (audit rights, incident reporting, sub-contracting controls, exit strategies). The Digital Omnibus harmonization (proposed November 2025) will streamline the overlapping AI Act / DORA / GDPR / NIS2 obligations.<sup>26</sup>

**Q6 — Is Vexorin a regulated entity?** A — No. Vexorin is not a registered broker-dealer, investment adviser, or exchange. Vexorin's outputs are evidence for customer analyst decisions; the customer is the regulated entity. Vexorin's architectural decision to position as a detection tool for institutional buyers (BRAIN Litepaper Section 12.5) means Vexorin's direct engagement with SEC FinHub or FINRA OFI is not a near-term GTM priority.

**Q7 — What happens if a customer is itself classified as a high-risk AI system deployer?** A — Vexorin's architecture is built to support customer compliance. The three-layer audit trail satisfies Article 12; the rule-based aggregation produces native explainability satisfying Article 13; the human-in-the-loop deployment model (analyst review for every high-conviction case) satisfies Article 14. Customers face a more credible compliance path with Vexorin than they do with black-box ML competitors.

**Q8 — How does Vexorin handle data licensing for the validated data marketplace?** A — Each data provider entering the marketplace executes a data-licensing agreement with Vexorin governing permitted uses, revenue share, customer-distribution rights, and provider warranties. The marketplace architecture (provider data validated against Vexorin's mesh, customer-acquired with proof of mesh integration) follows the data clean-room precedent established by platforms like Snowflake.

---

## 7. COMPLIANCE ROADMAP

---

The following commitments structure Vexorin's compliance buildout from pre-seed close through Series A.

**Days 1-30 post-pre-seed close.** Engage securities counsel for round structure and customer-contract templates. Engage data-protection counsel for GDPR and US state privacy law compliance review. Adopt standard institutional-vendor data processing agreement template. Establish OFAC screening process for customer onboarding.

**Days 30-90.** Deploy first three customers under Master Service Agreements with full DPA, sub-processor governance, audit rights, and incident-reporting clauses. Establish internal audit log retention policy. Complete EU AI Act general-purpose AI obligations compliance documentation.

**Days 90-180.** Engage regulatory-tracking infrastructure to monitor EU AI Act enforcement guidance, Digital Omnibus implementation, US state privacy law developments, and DORA technical standards. Begin SOC 2 Type I preparation for customers requiring formal certification. Position for SOC 2 Type II audit readiness during Series A window.

**Series A and beyond.** SOC 2 Type II certification. ISO 27001 certification. Specific certifications driven by customer requirements (FedRAMP for any US government customer; sector-specific certifications for European bank customers). Formal regulator engagement — not as a near-term GTM track, but as a credibility investment for Series B-stage customer expansion into the Tier-1 institutional segment.

---

## 8. WHAT THIS DOCUMENT IS NOT

---

This document is strategic analysis, not legal opinion. The compliance positions described here are informed by publicly-available regulatory documents and comparable-company precedents; they do not substitute for counsel-validated legal advice. Specific regulatory questions — particularly high-stakes ones (cross-border data flow for a specific customer, AI Act Annex III classification of a specific deployment configuration, GDPR Article 9 special-category data handling for a specific use case) — require counsel review before execution.

The document also is not a guarantee of compliance outcomes. Regulatory regimes evolve; enforcement priorities shift; comparable-company precedents do not bind regulators. Vexorin's compliance posture is built to be defensible, evidence-based, and aligned with industry-validated architectural choices — which is what survives a Tier-1 partner's data-room review and what positions the company to operate compliantly as it scales.

The document is what a senior partner reading Vexorin's data room expects to find: a substantive analysis of the regulatory landscape, anchored in comparable-company precedent, with the strategic and architectural choices that drive Vexorin's compliance posture explicitly stated.

---

## SOURCES

---

*Data Strategy & Compliance Position v1.0 · Vexorin.AI · May 2026*

---

1. Palantir Foundry technical specifications, deployment architecture, and compliance certifications. Source: invgate.com IT database entry on Palantir Foundry, accessed May 2026. <https://invgate.com/itdb/palantir-foundry>
2. Palantir Foundry technical specifications, deployment architecture, and compliance certifications. Source: invgate.com IT database entry on Palantir Foundry, accessed May 2026. <https://invgate.com/itdb/palantir-foundry>
3. EU AI Act (Regulation (EU) 2024/1689) — full text including Annex III high-risk categories and Article 9-43 obligations. Primary source: Official Journal of the European Union, June 13, 2024. [↗](#)
4. K&L Gates legal analysis: "EU and Luxembourg Update on the European Harmonised Rules on Artificial Intelligence — Recent Developments," January 20, 2026. <https://www.klgates.com/EU-and-Luxembourg-Update-on-the-European-Harmonised-Rules-on-Artificial-IntelligenceRecent-Developments-1-20-2026> [↗](#)
5. EU AI Act (Regulation (EU) 2024/1689) — full text including Annex III high-risk categories and Article 9-43 obligations. Primary source: Official Journal of the European Union, June 13, 2024. [↗](#)
6. K&L Gates legal analysis: "EU and Luxembourg Update on the European Harmonised Rules on Artificial Intelligence — Recent Developments," January 20, 2026. <https://www.klgates.com/EU-and-Luxembourg-Update-on-the-European-Harmonised-Rules-on-Artificial-IntelligenceRecent-Developments-1-20-2026> [↗](#)
7. Matproof analysis: "EU AI Act Compliance for FinTech: What You Need to Know," March 22, 2026. Includes the explicit statement that fraud detection AI is excluded from the high-risk credit scoring category. <https://matproof.com/blog/eu-ai-act-fintech-compliance> [↗](#)
8. Unit21 analysis: "EU AI Act 2026 FAQs: What Fraud and AML Teams Need to Know," May 2026. Confirms general-purpose AI obligations applied from August 2025. <https://www.unit21.ai/blog/eu-ai-act-2026-faqs-what-fraud-and-aml-teams-need-to-know> [↗](#)
9. Article 26 of the EU AI Act on deployer obligations of high-risk AI systems. Primary source: [artificialintelligenceact.eu/article/26/](https://artificialintelligenceact.eu/article/26/) [↗](#)
10. Dataminr Privacy Policy, accessed May 2026. <https://www.dataminr.com/privacy-policy/> [↗](#)
1. EquityZen company profile on Dataminr, including risk factors disclosing regulatory and compliance burden. <https://equityzen.com/company/dataminr/> [↗](#)
2. Dataminr Privacy Policy, accessed May 2026. <https://www.dataminr.com/privacy-policy/> [↗](#)
3. Digital Operational Resilience Act (Regulation (EU) 2022/2554), fully applicable from January 17, 2025. Primary source: Official Journal of the European Union. [↗](#)
4. Digital Omnibus proposal, European Commission, November 2025. Coverage in K&L Gates analysis cited in [^3]. [↗](#)

5. Matproof analysis: "EU AI Act Compliance for FinTech: What You Need to Know," March 22, 2026. Includes the explicit statement that fraud detection AI is excluded from the high-risk credit scoring category. <https://matproof.com/blog/eu-ai-act-fintech-compliance>
6. Vizologi business model analysis of Dataminr. <https://vizologi.com/business-strategy-canvas/dataminr-business-model-canvas/>
7. Dataminr Privacy Policy, accessed May 2026. <https://www.dataminr.com/privacy-policy/>
8. Vizologi business model analysis of Dataminr. <https://vizologi.com/business-strategy-canvas/dataminr-business-model-canvas/>
9. EquityZen company profile on Dataminr, including risk factors disclosing regulatory and compliance burden. <https://equityzen.com/company/dataminr/>
10. Dataminr Third Party Terms — X/Twitter Addendum. <https://www.dataminr.com/legal/thirdpartyterms/>
11. Palantir Foundry technical specifications, deployment architecture, and compliance certifications. Source: invgate.com IT database entry on Palantir Foundry, accessed May 2026. <https://invgate.com/itdb/palantir-foundry>
12. Matproof analysis: "EU AI Act Compliance for FinTech: What You Need to Know," March 22, 2026. Includes the explicit statement that fraud detection AI is excluded from the high-risk credit scoring category. <https://matproof.com/blog/eu-ai-act-fintech-compliance>
13. Palantir Foundry technical specifications, deployment architecture, and compliance certifications. Source: invgate.com IT database entry on Palantir Foundry, accessed May 2026. <https://invgate.com/itdb/palantir-foundry>
14. Dataminr Privacy Policy, accessed May 2026. <https://www.dataminr.com/privacy-policy/>
15. Vizologi business model analysis of Dataminr. <https://vizologi.com/business-strategy-canvas/dataminr-business-model-canvas/>
16. Digital Omnibus proposal, European Commission, November 2025. Coverage in K&L Gates analysis cited in [^3].