

vexorin faq investor objections

VEXORIN.AI — FAQ FOR KNOWN INVESTOR OBJECTIONS

Pre-Answered Diligence Questions for Tier-1 Investor Review · v1.0 · May 2026

Purpose. This document pre-answers the questions a Tier-1 venture partner is most likely to raise during diligence on Vexorin's pre-seed round. The objections are organized by category (technical, commercial, regulatory, team, market) and answered with the specific evidence, sourced reasoning, and architectural choices that defuse each one.

Audience. Senior partners at Founders Fund, Khosla, a16z American Dynamism, and the specialized fintech/regtech VCs who form the primary outreach target list. The document is written in the voice that survives a partner's diligence reading — direct, sourced, and honest about the limits where they exist.

Companion documents. This document references the BRAIN Litepaper, Master Reference, Prior Weight Matrix, 50-Case Manifold Pattern Analysis, Citations Audit, Data Strategy & Compliance Position, and the Risk Register (companion document in this data room). Where this FAQ summarizes a position, the companion documents carry the full underlying analysis.

1. EXECUTIVE NOTE

The most credible artifact a pre-seed company can place in a data room is an honest accounting of its own limits paired with a clear explanation of how those limits resolve over time. **The team that publishes a weaker honest position earns more trust than the team that publishes a stronger biased position, every time.** This document is calibrated to that principle.

Several of the questions below were originally surfaced in an internal audit conducted in May 2026, in which a senior reviewer adopting a Tier-1 IC posture stress-tested the deck and supporting materials and identified fifteen issues across technical, commercial, and strategic dimensions. Twelve of those issues are now closed in the deck and source documents; three are open commitments with stated mitigation plans. Both the closed and open issues are surfaced explicitly here.

2. TECHNICAL QUESTIONS

Q2.1 — You have one auditable real-time detection. Iran 2026. Everything else is retrospective reconstruction. How do you defend a category-defining claim from N=1?

The deck and the 50-Case Manifold Pattern Analysis are explicit that the library is **retrospective reconstruction, not a detection track record**. The Iran 2026 case is the only entry with a third-party-auditable real-time detection log. We surface this distinction directly in slide 7 panel 2 and in the Master Reference because the alternative — blurring the retrospective and real-time framings — is the kind of disclosure failure that converts a first call into a polite pass.

The structural defense to N=1 has three parts.

Part one — what the retrospective library actually proves. The signal weights underlying BRAIN are extracted from one hundred peer-reviewed academic citations and regulatory enforcement records, audited at publisher of record (Citations Audit v2.0: 90 verified clean, 7 metadata-corrected, 3 replaced with verified Tier-1 equivalents). The weights were not tuned to fit the cases. Applied retrospectively to fifty enforcement-confirmed manipulation events, the literature-derived signals produced seven geometrically distinct manipulation families with within-family cosine similarity ≥ 0.80 across all seven. The architecture's central falsifiable claim — that manipulation events leave coherent shapes across the manifold — survives the test against fifty cases the literature did not see during weight derivation.

Part two — the macro validation. SEC enforcement remedies grew from \$4.9B in FY2023 to \$8.2B in FY2024, approximately 67% year-over-year, the highest in the agency's history. Across financial-evidence categories where proof is structurally distributed across surfaces (money laundering being the clearest analog), interception rates for cases known to be occurring are estimated at or below 1% (FATF working estimates). The validation is structural: enforcement is scaling because the category is real and growing; detection infrastructure is the bottleneck; the visible portion of the category is the iceberg's tip.

Part three — what the round produces. This pre-seed funds the engineering buildout (4–5 month parallel-staffed effort, CTO-underwritten) that takes BRAIN aggregation from architectural specification to running code, brings three additional agents online (ECHO, ATLAS, CIPHER), and ships the Simulation Agent that opens the Validated Data Marketplace. By Series A, the system has moved from one auditable real-time detection (Iran) into deployed customer environments with Stream 1 (platform fees), Stream 2 (agent consumption), and Stream 3 (marketplace) all flowing. The N=1 problem resolves through deployment, not through additional retrospective claims.

Q2.2 — Your seven manipulation families were derived from the same fifty cases used to validate them. The cosine similarity numbers are in-sample. What’s the cross-validated number?

The point is correct and we surface it explicitly in Prior Weight Matrix v1.0.1 Part 6.1.

The Prior Weight Matrix derived seven family centroids from clustering the 50-case library, then validated the architecture by reporting within-family cosine similarity ≥ 0.80 and between-family ≤ 0.45 on five of six non-adjacent pairs. The test set equals the training set — this is methodologically a real limit, and the team that doesn’t surface it loses a partner who notices it on first read.

The forward commitment. Vexorin is running 5-fold cross-validation on the 50-case library. Methodology: hold out 10 cases per fold, derive 7 family centroids from the remaining 40, classify the held-out 10 against the derived centroids, repeat 5 times with different folds, report mean within-family and between-family cosine similarity on out-of-sample classifications.

The pre-registered acceptable thresholds. Within-family cosine similarity ≥ 0.75 (slightly relaxed from the in-sample 0.80 to allow honest out-of-sample degradation). Between-family cosine similarity ≤ 0.55 . Effect-size compression of 30–60% applied to weights in the cross-validation pass, so the reported numbers reflect compression-adjusted weights rather than raw literature magnitudes. **The thresholds are pre-registered before the test, not retrofitted to the result.**

If cross-validated numbers hold within these thresholds, the architecture’s central falsifiable claim survives the test against held-out data. If they weaken materially, the team will disclose the weaker numbers and adjust the family taxonomy and weight derivation accordingly. The output becomes Prior Weight Matrix v1.1, with cross-validated and compression-adjusted statistics replacing the in-sample uncompressed numbers as the primary reference.

The cross-validation is being run in the May 2026 post-prompt-finalization window. It is not a Series A commitment — it is a pre-seed-stage methodology refinement.

Q2.3 — The conviction equation multiplies by `pattern_factor [0.5, 2.0]` and `cascade_factor [1.0, 1.5]`. Where do those bounds come from?

The bounds are research-derived starting heuristics extracted from the 50-case library distribution and the 100 peer-reviewed citations. They are scaffolding, not free parameters. CALIBRATE refines them on customer-confirmed data over the first 60–90 days of deployment.

The honest framing: we chose conservative starting bounds informed by the library distribution. Customer-specific posteriors will adjust them within the first deployment window. The asymmetric range (penalty floor at 0.5×, amplification ceiling at 3×) is the structural defense against false positives — random non-cascading firing is penalized; clean cascades matching a canonical family are amplified.

This is the architectural reason BRAIN does not hallucinate. The bounds, the agent weights, and the family centroids are extracted from documented sources — peer-reviewed citations and enforcement-confirmed cases. The system is structurally constrained from recognizing patterns that were not first observed in documented ground truth. **The base image is real before the cascade can extend.**

The operational governance of these bounds is documented in BRAIN Litepaper Section 5 with the conviction-aggregation methodology. Hyperparameter governance — what’s tunable, what’s frozen, who approves changes — is part of the model governance regime in BRAIN Litepaper Section 12.3.

Q2.4 — 49,000 words of citation provenance is impressive. Where’s the code? Can I see one agent producing one alert against live data?

Scout and Carlton are production-ready. Scout deploys on-prem in 2–3 days with real-time coverage of every major social platform, 20,000+ news sites, 100+ languages. Carlton deploys alongside it for stock-price anomaly detection across the customer’s ticker universe. Customers interact with both like any institutional language-model deployment — query interface, continuous monitoring, full audit trail.

The architecture is **production-ready, pre-revenue at pre-seed by design**. Commercial GTM has not begun, and the round funds the conversion of validated demand (LOIs in the pipeline) into integrated deployments. The 49,000 words of

citation provenance establishes that the architecture is grounded; the production-ready agents establish that it runs; the round funds the connection between them.

For technical diligence, we offer a **Scout walkthrough** — data input pipeline, signal generator, YAML config, audit-trail emitter — as a 45-minute element of the first technical-diligence call. The partner sees the data fetcher, the signal generator, the configuration, and the audit-trail emission. Forty-five minutes of live code conversation closes the “is this real?” question more decisively than any document.

The clinching line. If we were generating signals at scale, we would not be at pre-seed. Live signal generation requires the data-feed licensing and pilot customer integration funded by this round.

Q2.5 — How does CALIBRATE actually update if the customer never sees a confirmed manipulation event in their first quarter?

This is a real architectural question and we surface the answer directly in BRAIN Litepaper Section 9 and in the Master Reference.

CALIBRATE refines on three feedback signals, not just confirmed events. Analyst-rejected alerts (false positives, plentiful — typically hundreds per customer per week). Analyst-confirmed alerts at any conviction tier (not just true manipulation). Regime-state classifications from AXIOM (continuous, derived from observed market structure).

The staged refinement timeline. Day 1: priors deployed. Week 4: first FPR refinement. Week 12: customer-specific FPR baseline. Year 1+: weight refinement on confirmed events as enforcement-confirmed cases accumulate.

The strategic framing. The slow accumulation of confirmed events is not a methodology limit — it is the category’s structural rarity, the same rarity Vexorin exists to address. If customers were detecting manipulation regularly with their existing tooling, they would not be evaluating Vexorin. The architecture matches the data rate; the data rate is the diagnosis.

Q2.6 — Effect-size compression of 30-60% on the foundation papers. What does that imply for the conviction scores BRAIN produces?

The compression is applied at deployment, not in customer environments after the fact. Customers do not encounter the optimistic published magnitudes; they encounter a calibrated baseline that reflects how these signals behave in the post-publication environment.

The Per-Agent Research Foundation Index documents the compression factors per agent — typically in the range [0.4, 0.7] depending on the signal class. Carry premium compression, insider-effect dilution by 10b5-1 plans, activist short-seller alpha decay, M&A run-up reduction by enforcement — these are well-documented post-publication effects in the academic-finance literature, and the foundation papers themselves often discuss them.

The cross-validation analysis described in Q2.2 is run with effect-size compression applied. The reported out-of-sample cosine similarity numbers in Prior Weight Matrix v1.1 will reflect compression-adjusted weights, not raw literature magnitudes. **This is the honest performance baseline.**

The disclosure converts a hidden risk into a visible architectural choice: a partner reading the data room sees that Vexorin has applied effect-size compression to every weight before deployment, which is the kind of disclosure-honest engineering choice that distinguishes a serious team from a pitch-optimized one.

Q2.7 — On-prem deployment vs marketplace coherence — how do independent developers ship code into closed customer environments?

The marketplace operates as a federated registry, not a hosted runtime. Independent developers publish verified, signed agent code or data sets to a centrally-verified registry. Customers browse the registry from inside their own on-prem deployment and pull verified modules into their environment when they want them. Vexorin handles registry governance — code review, signing, version pinning, security review, MCP-protocol compliance verification. Execution stays customer-side.

The architectural answer to the on-prem-vs-extensibility tension is not by shipping. The developer publishes verified code to the registry; the customer pulls. Customer pulls. Vexorin verifies. Nothing arrives without consent.

This is documented in BRAIN Litepaper Sections 20.1 through 20.4. Two extension layers operate on this customer-pull architecture — the Federated Registry for third-party signal modules, and the Validated Data Marketplace for third-party data sets validated through Vexorin’s Simulation Agent.

3. COMMERCIAL QUESTIONS

Q3.1 — Your TAM math is structurally optimistic. 3,700 institutions globally is a wider addressable than the realistic capture wedge. Walk me through the actual customer set.

The deck leads with **bottoms-up addressable, not top-down TAM**.

Initial GTM (Years 1-3): the underserved tier. Funds in the \$1-20B AUM range, regional and boutique asset managers, emerging-manager hedge funds, family offices with institutional ambition, second-tier sovereign-wealth allocators. They have the same recognition problem the Tier-1 funds have, without the budget for a 20-40-analyst surveillance team or the relationships to onboard fragmented alt-data. Sales cycle is shorter. Procurement is a fraction of a Tier-1 process. Several thousand institutions globally fit this profile. Target: 35-40 production customers in three years.

Long-term wedge (Years 3+): Tier-1 institutions. ~280 institutions globally with AUM > \$50B and active surveillance budgets > \$5M — the largest hedge funds (Citadel, Millennium, P72, Two Sigma, Renaissance), sovereign wealth funds (GIC, ADIA, Norges, KIA), and G-SIBs (JPM, GS, MS, BAML). These are the moat-defending customers, not the first-revenue customers.

Category context. The institutional surveillance and trade-compliance category is projected at \$9.31B by 2033 (17.25% CAGR — SNS Insider 2025). One supporting line, not the headline.

The Market Sizing Deep Dive (companion document, in build for the data room) extends this into specific named institutions, deployable surveillance budget estimates per sub-segment, and capture math sourced where possible to public 10-K and AUM disclosures.

Q3.2 — You propose a \$50K platform fee plus usage-based agent consumption plus a 15-20% data marketplace revenue share. Walk me through the buying journey.

The pricing model has three streams that diversify the income from day one.

Stream 1 — Platform fee (\$50K/year per customer). Predictable, recurring, the floor revenue line. Twelve active customers = \$600K ARR floor. This is what the customer pays for the deployed Vexorin runtime in their environment, regardless of usage volume.

Stream 2 — Agent consumption (usage-based). Scout at \$5K per million posts/articles/news items analyzed (\$0.005 per item). Carlton at \$20 per ticker per month. Future agents priced per usage profile relevant to their data class. Customer-value-aligned: customers who get more value pay more.

Stream 3 — Validated Data Marketplace. \$10K onboarding fee per data provider, paid at simulation analysis and deployment. 15-20% revenue share on data fees flowing through the validated marketplace to Vexorin’s customer base. Marketplace economics scale independently of customer count.

The buying journey is fast. Scout and Carlton deploy on-prem in 2-3 days. Customers interact with them like any institutional language-model deployment. There is no POC-to-pilot-to-production gauntlet — the agents are production-ready, deploy quickly, and the customer is in production from week one. The rest of the ecosystem integrates additively as it comes online.

The CFO underwrites the trajectory at every pitch. Forty years of institutional fundraising experience underwrites the projection — the Year 3 base case (~\$5.3M revenue at 12 active customers, 2.5% marketplace capture) is not a marketing forecast; it is a credentialed financial professional’s underwritten estimate.

Q3.3 — The data marketplace projection assumes 2.5% capture across 1,000 institutional data providers. How is that defensible?

We offer three scenarios rather than a single projection because the realistic capture rate is unknown at pre-seed:

Scenario	Capture	Providers	Onboarding (cumulative)	Revenue share ARR	Annual at maturity
Conservative	1%	10	\$100K	\$350K	~\$450K
Base (planning anchor)	2.5%	25	\$250K	\$875K	~\$1.1M
Strong	5%	50	\$500K	\$1.75M	~\$2.25M

The honest framing is “scenarios, not projections.” Three specific capture-rate scenarios let an investor underwrite a range of outcomes rather than a single number. The economics work at 1%; they work better at 5%. The deck doesn’t predict which level we hit — it shows the math is honest at any reasonable level.

Why 2.5% is defensible as the planning anchor. The Validated Data Marketplace solves a problem that does not have a competing solution: institutional alt-data providers cannot prove signal isolation in a way buyer research teams will sign off on. Each fund builds validation infrastructure badly, in-house, every time. Vexorin’s Simulation Agent produces a quantified mesh-integration validation artifact that providers carry into customer conversations. Two-and-a-

half percent capture on 1,000 providers — 25 onboarded providers across three years — is well below the natural ceiling of providers who would benefit from a Vexorin-branded validation artifact.

Q3.4 — Bloomberg is \$13B+ in revenue, 30,000+ employees, actively investing in surveillance and ESG/manipulation tooling. Why won't they catch up?

Bloomberg's surveillance architecture was built for a market that was single-surface. Vexorin complements Bloomberg by adding the multi-surface recognition layer rather than displacing the terminal. The architectural critique is structural — single-surface architectures cannot recognize multi-surface manipulation by construction — without claiming Bloomberg as a company is in decline.

The competitive landscape sits in seven categories (per the Competitive Landscape Analysis companion document). None of them produces what Vexorin produces. The category breakdown:

- Trade surveillance — exchange/regulator (NASDAQ SMARTS, Eventus, Trillium Surveyor, b-next): rules-based, order book / trade only.
- Bank-grade compliance surveillance (NICE Actimize, SteelEye): rules + ML, trade plus communications.
- Communications surveillance (Behavox, Smarsh): AI compliance, voice / chat / email.
- Fraud / anomaly ML (Featurespace, Quantexa): ML / entity resolution, single-domain.
- Crypto / on-chain (Solidus Labs, Chainalysis Markets, Elliptic, TRM Labs): rules + ML, on-chain only.
- In-house builds at Tier-1 funds (Citadel, Millennium, Two Sigma, Renaissance, P72): custom, fund-specific, not productized.
- Adjacent platforms (Palantir Foundry, general-purpose agent platforms): general-purpose, not domain-specific.

The white-space claim, sourced from literature. A search across the top ten lite papers in agentic prediction analysis and the top ten in manipulation detection plus geometric representation returned no architecture fusing both. **Agentic synthesis plus geometric manifold representation has not been published.** The white space is in the literature itself, verifiable on arXiv and SSRN.

Q3.5 — Your most likely competitor is not another vendor — it's a Tier-1 hedge fund building this internally. How do you respond?

Citadel, Millennium, Two Sigma, Renaissance, and P72 do have internal teams capable of building similar systems. They are not productizing for two structural reasons.

One — building an architecture that survives multi-customer deployment is a different engineering problem than building one for internal use. An internal team optimizes for one fund's data infrastructure, one fund's regulatory regime, one fund's operating culture. A productized system has to handle data heterogeneity, multi-jurisdictional compliance, customer-side governance, and an upgrade path that survives across customer environments. The architectural investment to productize is materially larger than the investment to deploy internally.

Two — the strategic incentive to productize is weak. A Tier-1 fund's competitive advantage from a recognition substrate is enhanced by *not* productizing. Sharing the architecture — even commercially — reduces the relative advantage. The internal-build strategy is a stable equilibrium for the Tier-1 funds; productization is a strategic decision they are unlikely to make.

Vexorin's edge against the in-house build threat. Purpose-built for cross-institution deployment from day one. On-prem from day one (which Tier-1 funds expect for their own infrastructure regardless). Sourced and audited foundation that an internal build cannot replicate without doubling its team size to do the citation provenance work. Falsifiability record (cosine similarity tests, pre-registered thresholds) that compounds across customer deployments. **The Tier-1 internal build is not Vexorin's competitor — it is Vexorin's eventual customer once the product matures into the long-term wedge.**

Q3.6 — How does Vexorin handle the procurement gauntlet at institutional buyers? Isn't the sales cycle 9-18 months for this category?

Vexorin's primary GTM is the underserved tier — mid-AUM funds, regional and boutique asset managers, emerging managers, family offices. This segment has materially shorter procurement cycles than Tier-1 institutions. They have the same recognition problem without the procurement gauntlet that kills sub-\$10M startups.

Three founder-owned distribution paths compound for these customers. Amir Zada's MD-level history at Renaissance Capital, Exotix Ltd, and Global Frontera Markets, plus his current G3 Good Governance Group role and TLG Capital advisory, opens fund procurement conversations. Peter (CTO) owns language-model term optimization — Vexorin's appearance in AI-search and LLM-query results, which the 2026 institutional buyer increasingly relies on. Alistair Ewart's twenty years of understanding virality algorithms and how to optimize them produces organic distribution within the networked conversations on Reddit, Twitter/X, and financial-media-adjacent communities where institutional

buyers are increasingly informed.

The three paths cannot be replicated by a generalist team without recruiting three specialized founders. They compound: Amir's direct sales convert a fund; Peter's LLM term optimization makes that fund's research analyst find Vexorin when they query Claude or ChatGPT; Alistair's organic distribution makes the same analyst already familiar with Vexorin from networked conversations before the procurement conversation begins. Three independent paths to the same buyer.

The Customer Acquisition Strategy companion document expands this into specific target lists, outreach motion per tier, and conversion metrics.

4. REGULATORY QUESTIONS

Q4.1 — Where does Vexorin sit under the EU AI Act?

Vexorin is not classified as a high-risk AI system under EU AI Act Annex III. Annex III's financial-services high-risk categories are credit scoring (point 5b), insurance pricing (point 5c), and biometric identification (point 1). Vexorin's manipulation detection use case sits closer to fraud detection, which is explicitly excluded from the high-risk credit-scoring category.

The Act's general-purpose AI obligations applied from August 2025 — transparency requirements, technical documentation, copyright compliance — do apply to Vexorin and the platform is built to satisfy them through its rule-based aggregation architecture and three-layer audit trail.

The strategic positioning advantage. Customers deploying Vexorin to support their own high-risk AI obligations face a more credible compliance path than they would with black-box ML competitors. BRAIN's three-layer auditable provenance directly satisfies Article 12 audit-trail requirements; the rule-based aggregation produces native explainability satisfying Article 13; the human-in-the-loop deployment satisfies Article 14.

The Data Strategy & Compliance Position document (companion to this FAQ) extends this analysis with primary-source citations and the comparable-company precedent (Dataminr for public-data ingestion, Palantir Foundry for on-prem deployment).

Q4.2 — How does Vexorin handle GDPR?

Architectural choices materially constrain GDPR exposure.

Public-data-only ingestion at the architecture level. Scout's coverage is restricted to publicly-available content by architectural design. Public social media posts contain personal data, but processing of public content is broadly permitted under GDPR's "legitimate interest" basis (Article 6(1)(f)). Dataminr operates under this same architecture and has scaled to global Tier-1 financial-institutional customers under GDPR.

Customer data handling is governed by the customer. Customer order flow, portfolio data, and internal communications are processed on-prem under the customer's own GDPR controls. Vexorin acts as a data processor (or sub-processor) under Article 28; the relationship is governed by Data Processing Agreements with audit rights, sub-processor approval, and incident reporting.

Right-to-be-forgotten architecture. Scout's ingestion respects platform-level deletion: when a social media user deletes a public post, the post is removed from Scout's monitoring scope on the next ingestion cycle.

Q4.3 — Are you a regulated entity? Do you need SEC, FINRA, or FCA authorization?

No. Vexorin is not a registered broker-dealer, investment adviser, or exchange. Vexorin's outputs are evidence for customer analyst decisions; the customer is the regulated entity.

Vexorin's architectural decision to position as a detection tool for institutional buyers — not as a regulator-targeted product — means direct regulatory engagement with SEC FinHub or FINRA OFI is not a near-term GTM priority. The architecture is built to **support** customer regulatory workflows; Vexorin is not itself a regulated entity.

This is documented in BRAIN Litepaper Section 12.5.

Q4.4 — What about DORA?

Vexorin's exposure under DORA is as a third-party ICT service provider to financial-entity customers. DORA's third-party ICT risk management framework (Articles 28-30) requires customers to manage their critical ICT third-party providers under contractual frameworks that include audit rights, incident reporting, sub-contracting controls, and exit strategies.

Vexorin's customer contracts will include these clauses standard. The Digital Omnibus proposal (European Commission, November 2025) simplifies the overlapping AI Act / GDPR / NIS2 / DORA / Data Act framework with a single incident-

reporting point and aligned breach-notification thresholds. Vexorin will track Digital Omnibus implementation and integrate the harmonized framework as the regime stabilizes.

Q4.5 — How do you handle data licensing for the social and news data Scout ingests?

Scout's coverage is structurally restricted to publicly-available content. The licensing posture mirrors Dataminr's, which has scaled to Tier-1 financial-institutional customers under EU and US regulatory regimes.

For platform-specific content (e.g., X/Twitter, Reddit, specific news platforms), licensing relationships are negotiated platform-by-platform under the relevant Developer Terms or commercial licensing agreements. Dataminr's Third Party Terms (publicly disclosed) document the standard architecture: the platform licenses content to Dataminr; Dataminr licenses derived analytical product to its customers; the platform reserves third-party-beneficiary enforcement rights.

Vexorin's data feed licensing budget (\$90K in the use of funds allocation) covers the platform-specific licensing for Scout's 100-language coverage at customer scale.

5. TEAM QUESTIONS

Q5.1 — Walk me through who on this team has sold institutional software to Tier-1 hedge funds before, who has built compliance infrastructure at scale, and who has a relationship with a regulator.

The team's institutional-finance pedigree is unusually deep for the category.

Alistair Ewart, Co-CEO. Visionary founder of KOAT.AI — a government-deployed online social manipulation detection tool. Twenty years of understanding virality algorithms and how to optimize them. Scout, and the social-layer recognition substrate of the manifold, are the productionized continuation of this work. The institutional knowledge transfer with the founder.

Amir Zada, Co-CEO. Recruited into British intelligence from high school. Director and Co-Head, New York Office, G3 Good Governance Group (intelligence and cyber defense). Prior Managing Director at Renaissance Capital and Exotix Ltd, and Senior Managing Director at Global Frontera Markets — frontier-markets investment banking, where institutional manipulation is most prevalent and least surveilled. MSc Management, London School of Economics. MSc International Management, HEC Paris. Member, Board of Advisors at TLG Capital. **The Renaissance Capital and Exotix MD-level history is the direct answer to “who has sold to Tier-1 institutions before.”** It has been done, at MD level, across multiple firms, in markets where the buying journey for institutional surveillance is most demanding.

CTO. Multiple prior exits. Operating credibility for the engineering buildout. (Detail in companion documents.)

CFO. Forty years of institutional fundraising experience. Validates sales trajectory at every pitch. (Detail in companion documents.)

Engineering Advisor. Director of Engineering at Kraken. Real-time financial-infrastructure operating credibility at scale; surveillance, microstructure, and on-chain mechanics intersect in his daily work. **Holds a direct connect into Founders Fund — the warm-intro path for Tier-1 outreach.**

The team-as-thesis frame: the architecture follows from who they are. Each founder's institutional history corresponds to a layer of the architecture. Replication requires going back twenty years.

Q5.2 — Why is this the team that builds this specific architecture?

The architecture is the artifact of the team's specific operating history, and only this assembly of credentialed humans could have produced it.

KOAT.AI's government-deployed social-manipulation detection tool is the substrate Scout is built on. Alistair's twenty years inside virality-algorithm operating expertise produced the institutional knowledge that Scout's social-layer recognition substrate inherits. **The IP and operating history transfer with the founder; replication does not.**

Amir's British intelligence fieldwork plus MD-level frontier-markets investment banking pedigree produced the institutional capital-markets and intelligence-coordination understanding that the manifold's broader surfaces (Disclosure, Intent, Positioning, Flow) are designed against. Frontier markets are where institutional manipulation is most prevalent and least surveilled, which means Amir's operational understanding of the manipulation landscape is materially deeper than a generalist banker's.

The Engineering Advisor's role at Kraken supplies real-time financial-infrastructure operating credibility for the integration phase. The CTO's prior exits supply the engineering execution credibility for the parallel-staffed build. The CFO's forty years of fundraising underwrites the financial trajectory.

Each credential corresponds to a specific layer of the architecture, and the layers compose. This is the

strongest single competitive moat the company has — and it is the moat investors recognize and underwrite, because it cannot be replicated by a generalist team without going back twenty years.

Q5.3 — Is the team financially aligned for the long term? Vesting, IP assignment, founders' agreements?

Cap table, vesting terms, founders' agreements, and IP assignments are documented separately in the data room (CFO-owned legal documentation set). All four founders are subject to standard four-year vesting with one-year cliff. IP assignment agreements have been executed for KOAT.AI's substrate transfer to Vexorin and for all founder-contributed engineering work. Advisor agreements are in place for the Engineering Advisor and additional named advisors.

6. MARKET QUESTIONS

Q6.1 — Surveillance and regtech aren't categories Founders Fund / a16z / Sequoia hunt at pre-seed. How is this round structured to land?

The fundraising strategy reflects the structural constraint.

Primary path — specialized fintech/regtech pre-seed VCs. Susa Ventures, Bain Capital Ventures' fintech-specific partners, Inspired Capital, NEA's fintech focus, Nasdaq Ventures, Citi Ventures, MGVC, certain partners at Lerer Hippeau and Initialized. These funds have partners with direct surveillance and compliance domain knowledge.

Active warm-intro path — Founders Fund via Engineering Advisor's direct connect. Trae Stephens is the obvious thesis fit. The Anduril-adjacent intelligence-community thesis maps directly to the team's KOAT.AI government deployment and British intelligence fieldwork. The warm intro path converts at materially higher rates than cold outreach.

Adjacent paths. a16z American Dynamism (Boyle/Ulevitch — the intelligence-community thesis applies); a16z fintech infrastructure (Rampell/Strange — the institutional-buyer journey applies); Sequoia (selective).

The cap-table-construction frame. This pre-seed builds the foundation cap table. Series A brings the Tier-1 lead and the brand validation that compounds through Series B. The pre-seed cohort is not a placeholder for "real" investors who arrive later — it is the foundation layer of a deliberately constructed cap table where each round adds capability the company needs at that stage.

Q6.2 — What's the realistic path to Series A? On what timeline and at what valuation?

Series A at Month 10–12, \$8–12M valuation. The architecture moves from one auditable real-time detection (Iran) into deployed customer environments with all three revenue streams flowing.

The six-month execution arc (per slide 13 of the deck): Today — production-ready Scout and Carlton, LOIs in pipeline, CFO-underwritten projections. Months 1–2 — first three customers deploy in 2–3 days, Stream 1 and Stream 2 begin flowing, engineering parallel streams begin. Months 3–4 — first cohort fully operational, cross-validation analytical pass complete, customer count expanding. Month 5 — BRAIN aggregation engine operational, three additional agents integrated. Month 6 — Simulation Agent operational, Validated Data Marketplace opens, first data providers onboarded, Stream 3 begins flowing.

Series A conversation goes live at Month 6+, anchored on revenue traction across three streams plus the deployed-ecosystem demonstration. The Tier-1 lead at Series A brings the brand validation that compounds through Series B.

Q6.3 — What's the worst-case scenario and how do you respond to it?

The worst case is realistic and the team has thought about it.

If first-customer deployment takes longer than projected, Stream 1 and Stream 2 ramp slower than the base case. The \$750K runway extends ~24 months post-close, which provides material buffer against deployment delays. The CFO underwrites the trajectory at every pitch and tracks against the planning anchor monthly.

If the Simulation Agent build takes 6 months instead of 4–5, Stream 3 launches in Month 8 rather than Month 6. The marketplace ramps slower; Year 1 revenue is more heavily front-loaded on Streams 1 and 2; Year 3 base case stays at ~\$5.3M because the marketplace has time to mature.

If cross-validation produces materially weaker out-of-sample numbers than the in-sample reference, the team discloses the weaker numbers and adjusts the family taxonomy. Trust compounds; the architectural argument shifts to the structural validity of the manifold approach rather than to specific cosine numbers.

If a Tier-1 customer in-house build emerges as productized competition (low probability per Q3.5), Vexorin's purpose-built-for-cross-institution-deployment architecture and citation-audited foundation provide defensive moat that internal

builds cannot replicate without doubling team size.

The Risk Register companion document organizes this analysis into a structured exposure framework with mitigation status and owner per category.

7. CLOSING NOTES

This FAQ covers the most likely diligence questions. The data room includes additional companion documents that extend the analysis:

- **Risk Register** — categorized exposure analysis with mitigation status
- **Competitive Landscape Analysis** — per-vendor profiles and white-space defense
- **IP & Architecture Defensibility** — moat slide expanded with sourced provenance
- **Customer Acquisition Strategy** — three-tiered GTM with target lists and conversion metrics
- **Data Strategy & Compliance Position** — regulatory landscape with comparable-company precedent
- **Market Sizing Deep Dive** — bottoms-up customer analysis (in build)
- **Six-Month Operating Plan** — weekly milestones (in build)
- **Financial Model with Sensitivity Analysis** — three-stream revenue framework (CFO-owned, in build)
- **Use of Funds Detailed Allocation** — line-by-line breakdown (CFO-owned, in build)

Questions not covered in this FAQ are typically routed to the relevant companion document. Questions that arise during diligence and are not pre-answered here will be answered directly by the founder team — Alistair, Amir, or the CTO depending on subject — within 48 hours of the question.

The team welcomes diligence at depth. **The architecture is testable. The team is credentialed. The path is documented. The work is integration.**
