vexorin master reference

VEXORIN.AI — MASTER REFERENCE

Canonical Strategic Reference · v2.0 · May 2026

> ***Purpose.*** *This document is the single canonical strategic and commercial reference for Vexorin.AI. The BRAIN Litepaper carries the architectural and methodological depth; this document carries the company's identity, team, market positioning, GTM, fundraising structure, and locked language. Together they constitute the complete operational reference for the company.*
>
> ***When to use this document.*** *When a new advisor, prospective investor, or partner needs to understand Vexorin in its entirety in one read. When generating any deck, pitch, email, or external communication. When checking that strategic claims across documents are consistent with the company's actual positioning.*
>
> ***Version note.*** *This document supersedes any prior Master Reference. The language, GTM structure, pricing tiers, fundraising terms, team credentials, and competitive framing reflect the company's current state as of May 2026 and the strategic decisions made through the deck and audit work of April–May 2026.*

---

1. WHAT VEXORIN IS

1.1 The one-sentence answer

Vexorin is the recognition substrate for institutional capital — a multi-agent architecture that recognizes coordinated market manipulation by reading the geometric shape it leaves across six evidence surfaces, before broader market recognition catches up.

1.2 The locked vocabulary

Vexorin recognizes; it does not predict. It surfaces indicators ahead of broader recognition. The system reads the manifold; it does not detect events. Manipulation deforms the manifold; the architecture recognizes the deformation as a signature.

Always use: recognition, recognize, surface, indicate, deformation, signature, shape, manifold, cascade, geometry, substrate.

Never use: predict, predicted, detected, forecast, alpha generation, trade signal, alert (as a verb describing system behavior), AI prediction, machine learning forecasting.

1.3 What Vexorin is not

Vexorin is not a trade signal generator. Vexorin is not a regulator. Vexorin is not a replacement for human analysts. Vexorin is not a black box — its conviction aggregation is deliberately rule-based and traceable. Vexorin's day-one performance is research-derived, not customer-tailored, and the deck and all customer-facing materials reflect this honestly.

1.4 The brand identity

Brand colors — VOID #0A0A0F (near-black background), DEPTH #0F1219 (panel background), STEEL #1A2B4A (accent dark), SIGNAL #00E5FF (primary highlight), ALERT #FFB800 (secondary highlight), THREAT #FF2D55 (warning highlight). Type — restrained, mono-treatment for headers, sans-serif body. Visual posture — geometric, not decorative; type-driven over icon-driven; standard background throughout; one chart maximum per deck.

---

2. THE COMPANY THESIS

2.1 The single thesis

The institutional surveillance category is structurally broken because it observes single surfaces, while modern coordinated manipulation distributes its evidence across multiple surfaces by design. The cases prosecuted are by definition the cases where single-surface evidence was sufficient. The harder, larger, faster-growing category is structurally invisible to current surveillance infrastructure. Vexorin closes this gap by reading all six surfaces as one geometry.

2.2 Why now

SEC financial remedies obtained reached $8.2B in FY2024, up from $4.9B in FY2023 — approximately 67% year-over-year growth, the highest in the agency's history. Enforcement is scaling to meet an expanding category, but enforcement depends on detection. Detection infrastructure is structurally limited to single-surface observation, which means enforcement is being constrained by what its tooling can support. Recognition is the bottleneck. Prosecution is downstream of recognition. The visible portion of the category is growing fast; the invisible portion — manipulation events that cross surfaces by construction — is structurally larger and currently uninstrumented.

2.3 The iceberg framing

Across financial-evidence categories where proof is structurally distributed across surfaces (money laundering being the clearest analog), interception rates for cases known to be occurring are estimated at or below 1% (FATF working estimates). There is no architectural reason multi-surface manipulation would be different, and several structural reasons to believe it is the same or worse. The category's true size is not measurable from prosecuted cases; it is materially larger, growing materially faster, and currently invisible to the surveillance infrastructure deployed against it.

---

3. THE TEAM (TEAM-AS-THESIS)

3.1 The frame

The architecture follows from who they are. Each founder's institutional history corresponds to a layer of the architecture, and the layers compose. The team is not a competent collection of operators with credible advisors — it is the only assembly of people who could have built this specific architecture. Replication requires going back twenty years.

3.2 Founders

Alistair Ewart, Co-CEO. Visionary founder of KOAT.AI — a government-deployed online social manipulation detection tool. Twenty years of understanding virality algorithms and how to optimize them. Scout, and the social-layer recognition substrate of the manifold, are the productionized continuation of this work. The institutional knowledge transfer with the founder.

Amir Zada, Co-CEO. Recruited into British intelligence from high school. Director and Co-Head, New York Office, G3 Good Governance Group (intelligence and cyber defense). Prior Managing Director at Renaissance Capital and Exotix Ltd, and Senior Managing Director at Global Frontera Markets — frontier-markets investment banking, where institutional manipulation is most prevalent and least surveilled. MSc Management, London School of Economics. MSc International Management, HEC Paris. Member, Board of Advisors at TLG Capital. The institutional capital-markets and intelligence-coordination layers of the manifold follow directly from this fieldwork.

CTO. Multiple prior exits. Operating credibility for the integration phase. (Specific name and exit history to be added when this document is shared externally.)

CFO. Forty years of institutional fundraising experience. Validates sales trajectory at every pitch — the financial projections are not marketing forecasts; they are a credentialed financial professional's underwritten estimates. (Specific name and notable raises to be added when this document is shared externally.)

3.3 Advisors

Engineering Advisor — Director of Engineering at Kraken. Real-time financial-infrastructure operating credibility at scale. Surveillance, microstructure, and on-chain mechanics intersect in his daily work. Holds a direct connect into Founders Fund, which is the warm-intro path for Tier-1 outreach.

Chief Strategy Advisor. (Credentials to be reconciled with source documents and added.)

Architecture Advisor. (Credentials to be reconciled with source documents and added.)

3.4 The team-as-thesis frame for outreach

Two operators with two decades inside virality-algorithm and intelligence-community signal work, with a government-deployed manipulation detection tool and an British intelligence recruitment in their backgrounds, are now applying that infrastructure thesis to institutional capital markets. The architecture follows from the team's specific operating history. Recognition is the new alpha, and we are the only team with the institutional history to build it.

4. THE ARCHITECTURE (SHORT REFERENCE)

> *Full architectural detail lives in the BRAIN Litepaper. This section is the strategic-pitch summary of what BRAIN is and how it works, calibrated for a non-technical reader.*

4.1 The substrate

Vexorin is built on a six-surface manifold — Context, Disclosure, Intent, Positioning, Flow, Structure. The six surfaces fire in cascade order during a manipulation event: outer-context narrative formation precedes disclosure activity precedes intent activity precedes positioning shifts precedes flow imbalance precedes structural fingerprint. Random correlated firings do not produce cascade structure. Manipulation events do, in documented temporal order. The cascade is the shape.

4.2 The agents

Ten agents read the surfaces — SCOUT (14.7%, Context, social-layer), ECHO (17.6%, Flow, options/vol/skew), ATLAS (27.4%, Disclosure, financial sentiment + microstructure), CIPHER (18.0%, Intent, insider footprint), NEXUS (9.8%, Positioning, cross-asset), HERALD (6.4%, Disclosure, event study), ARGUS (4.6%, Context, geopolitical), MAGNUS (5.4%, Context, monetary policy), ORION (3.8%, Context, cross-currency), AXIOM (2.0%, Structure). Reserved capacity 4.4% for v1.0 expansion. AXIOM is the regime classifier — broadcasting SYSTEM_QUIET / ELEVATED / STRESSED / CRISIS — not a primary signal generator; it modifies every other agent's confidence by ±0.10 to ±0.20.

4.3 The math

BRAIN's manifold conviction is `base × correlation_factor × pattern_factor × cascade_factor`, with pattern_factor bounded [0.5, 2.0] and cascade_factor bounded [1.0, 1.5]. Random non-cascading firing is penalized to 0.5× the naive baseline; clean cascades matching a canonical family are amplified to 3×. The asymmetry is the structural defense against false positives.

4.4 Why BRAIN does not hallucinate

The bounds, the agent weights, and the family centroids are all extracted from documented sources — peer-reviewed citations and enforcement-confirmed cases. They are scaffolding, not free parameters. The system is structurally constrained from recognizing patterns that were not first observed in documented ground truth. The base image is real before the cascade can extend.

4.5 The seven manipulation families

Cases cluster into seven canonical families — A Social Squeeze, B Classical Insider, C Microstructural Spoofing, D Pump-and-Dump, E Geopolitical/Commodity, F Oracle/Governance, G Accounting Fraud. Each family passes within-family cosine similarity ≥ 0.80; between-family cosine similarity ≤ 0.45 on five of six non-adjacent pairs. The architecture is testable, and it has passed its own test.

4.6 Two extension layers

Federated registry — independent developers publish MCP-compatible signal modules to a centrally-verified registry; customers pull verified modules into their on-prem deployment when they want them. **Validated data marketplace** — third-party data providers submit data sets to Vexorin's mesh, simulation runs validate signal contribution, validated sets earn placement and customers acquire them with proof of mesh integration attached. Both layers customer-pulled. Vexorin verifies. Execution stays customer-side. Nothing arrives without consent.

5. VALIDATION POSTURE

5.1 What has been validated

The signal weights underlying BRAIN are extracted from one hundred peer-reviewed academic citations and regulatory enforcement records, audited at publisher of record (90 verified clean, 7 metadata-corrected, 3 replaced with verified Tier-1 equivalents). The weights were not tuned to fit cases. Applied retrospectively to fifty enforcement-confirmed manipulation events filed by SEC, DOJ, CFTC, and FINRA between 2022 and early 2026, the literature-derived signals produced seven geometrically distinct manipulation families. Forty-seven of fifty cases deformed two or more surfaces; thirty-one of fifty deformed three or more.

5.2 The honest limits

The 50-case library is **retrospective reconstruction**, not a detection track record. The Iran 2026 case is the only entry with a third-party-auditable real-time detection log: BRAIN Conviction 76 issued January 4–14, 2026 while WTI traded $63–65; kinetic event February 28; WTI peaked $119.50 on March 9; forty-five days ahead of broader recognition. Cosine similarity 0.91 against the Family E centroid, fitting tighter than four of the five other Family E members.

5.3 What is being validated next

Out-of-sample 5-fold cross-validation on the seven families, with effect-size compression of 30–60% applied to the weights at the same time. Pre-registered acceptable thresholds: within-family cosine ≥ 0.75, between-family cosine ≤ 0.55. The cross-validated and compression-adjusted numbers will replace the in-sample uncompressed numbers in customer-facing materials when the analytical pass completes.

5.4 The v0.1 priors framing

The current weight matrix is the day-one deployment vector — a research-derived prior, not a validated posterior. CALIBRATE produces customer-specific posteriors over time. The 60–90 day window covers customer-specific FPR recalibration on triage outcomes (analyst-rejected and analyst-confirmed alerts, which are abundant). Full posterior weight refinement on confirmed manipulation events runs over 12+ months because confirmed events are rare — and that rarity is the structural condition Vexorin exists to address.

---

6. THE MARKET (BOTTOMS-UP, NOT TOP-DOWN)

6.1 Initial GTM (Years 1–3): the underserved tier

Funds in the $1–20B AUM range, regional and boutique asset managers, emerging-manager hedge funds, family offices with institutional ambition, and second-tier sovereign-wealth allocators. They have the same recognition problem the Tier-1 funds have, without the budget for a 20–40-analyst surveillance team or the relationships to onboard fragmented alt-data. Sales cycle is shorter. Procurement is a fraction of a Tier-1 process. Tape is meaningfully lower. Several thousand institutions globally fit this profile. Target: 35–40 production customers in three years.

6.2 Long-term wedge (Years 3+): Tier-1 institutions

~280 institutions globally with AUM > $50B and active surveillance budgets > $5M. The largest hedge funds (Citadel, Millennium, P72, Two Sigma, Renaissance), sovereign wealth funds (GIC, ADIA, Norges, KIA), and G-SIBs (JPM, GS, MS, BAML). These are the moat-defending customers, not the first-revenue customers. The bridge from initial GTM to this tier is the customer-validated posteriors compounding from the first cohort and the architecture's track record growing through enforcement-confirmed events.

6.3 Category context

The institutional surveillance and trade-compliance category is projected at $9.31B by 2033 (17.25% CAGR — SNS Insider 2025). One supporting line, not the headline.

---

7. PRICING (THREE STREAMS, DIVERSIFIED FROM DAY ONE)

7.1 The principle

Vexorin's commercial model is built on three distinct revenue streams. Customers pay for a deployed ecosystem, for the agents they activate against their data, and for validated marketplace integrations. The architecture supports diversified income from day one, not a single-line projection.

7.2 Stream 1 — Platform fee

$50,000/year per customer for the on-prem ecosystem deployment. Predictable, recurring, the floor revenue line. Twelve active customers = $600K ARR floor.

7.3 Stream 2 — Agent consumption (usage-based)

- **Scout** — $5,000 per million posts/articles/news items analyzed ($0.005 per item). Real-time coverage of every major social platform, 20,000+ news sites, 100+ languages.
- **Carlton** — $20 per ticker per month. Stock-price anomaly detection across the customer's ticker universe.

- **Future agents** (ECHO, ATLAS, CIPHER, NEXUS, HERALD, ARGUS, MAGNUS, ORION, AXIOM) — each priced per usage profile relevant to its data class. New internal agents cost approximately $15,000 to develop with a 2–4 week build cycle, so the catalog grows continuously.

Customer-value-aligned: customers who get more value (more tickers, more conversation volume, more agents activated) pay more.

7.4 Stream 3 — Validated Data Marketplace economics

- **Onboarding fee** — $10,000 per data provider, paid at simulation analysis and deployment.
- **Revenue share** — 15–20% of the data provider's fees on data flowing through the validated marketplace to Vexorin's customer base.

Marketplace economics scale independently of customer count.

7.5 Three-scenario marketplace projection

Across an addressable universe of approximately 1,000 institutional data providers globally:

Scenario	Capture	Providers	Onboarding (cumulative)	Revenue share ARR	Annual at maturity
Conservative	1%	10	$100K	$350K	~$450K
Base (planning anchor)	**2.5%**	**25**	**$250K**	**$875K**	**~$1.1M**
Strong	5%	50	$500K	$1.75M	~$2.25M

Revenue share assumes ~$200K/year in data fees per onboarded provider at 17.5% share rate. CFO underwrites the base case as the planning anchor.

7.6 Composite Year 3 revenue (base case)

12 active customers, 2.5% marketplace capture: $600K platform fees + ~$3.6M agent consumption + $875K marketplace revenue share + ~$250K cumulative onboarding fees = **~$5.3M revenue Year 3**. CFO validates the trajectory at every pitch — forty years of institutional fundraising experience underwrites the projection.

7.7 Deployment timeline

Scout and Carlton are production-ready and deploy on-prem in 2–3 days. Customers interact with deployed agents like a language model — query interface, continuous monitoring, audit trail. The rest of the ecosystem integrates additively as it comes online; customers do not wait for "the platform" to be complete because they are in production from week one.

7B. GO-TO-MARKET (THREE FOUNDER-OWNED PATHS)

7B.1 The principle

The architecture follows from who they are. So does the distribution. Each founder's institutional history corresponds to a distribution path the team can execute that a generalist team cannot replicate. The three paths compound rather than substitute — three independent paths to the same buyer.

7B.2 Tier 1 — Direct fund relationships (Amir)

Amir's MD-level history at Renaissance Capital, Exotix, and Global Frontera Markets, plus his current G3 Good Governance Group role and TLG Capital advisory, opens fund procurement conversations. Direct enterprise sales motion: credibility-driven, relationship-driven, slow-but-high-value. Mid-tier funds, family offices, sovereign-wealth secondary tiers. Converts warm relationships into LOIs and LOIs into deployed Scout+Carlton instances.

7B.3 Tier 2 — Language model term optimization (Peter)

Peter's CTO credibility and agentic engineering depth makes him the right founder to own technical positioning in AI search and LLM term spaces. The 2026 institutional buyer increasingly discovers products through AI search rather than traditional Google search — they query Claude or ChatGPT for surveillance products, integration patterns, technical comparisons. Peter owns Vexorin's appearance in those answers through structured documentation, technical commons surfacing, and LLM-aware product positioning. A 2026-native distribution strategy that did not exist three years ago.

7B.4 Tier 3 — Organic distribution through networked conversation platforms (Alistair)

Alistair's twenty years of understanding virality algorithms and how to optimize them translate directly into Vexorin's organic distribution. Surfacing Vexorin within the networked conversations on Reddit, Twitter/X, and the financial-media-adjacent communities where institutional buyers are increasingly informed. Founder-as-distribution-channel in the most literal sense: the founder who understands the substrate Scout reads is the founder who can use that substrate to surface Vexorin within it.

7B.5 Why the three-tiered GTM is itself a moat

Each path exploits a credential specific to the founder owning it. The three paths cannot be replicated by a generalist team without recruiting three specialized founders. They compound: Amir's direct sales convert a fund; Peter's LLM term optimization makes that fund's research analyst find Vexorin when they query Claude; Alistair's organic distribution makes the same analyst already familiar with Vexorin from networked conversations before the procurement conversation begins. *The architecture follows from who they are. So does the distribution.*

---

8. THE MOAT

Five components:

1. **Falsifiable architecture.** Pre-registered criteria define when the system would be declared incorrect. Within-family cosine similarity must stay ≥ 0.80; between-family cosine ≤ 0.45 on non-adjacent pairs. The architecture is testable, and it has passed its own test. The falsifiability record cannot be backdated.
2. **49,000-word research foundation.** Ten dedicated agent research documents, six peer-reviewed foundation papers each, all 100 Tier-1 citations independently verified at publisher of record. Effect-size compression of 30–60% is applied to every weight at deployment — customers do not encounter the optimistic published magnitudes.
3. **CALIBRATE walk-forward.** Refines continuously on triage outcomes for customer-specific FPR recalibration in 60–90 days; weight refinement on confirmed manipulation events runs over 12+ months. The slow accumulation of confirmed events is not a methodology limit — it is the category's structural rarity, the same rarity Vexorin exists to address.
4. **50-case validated library at deployment.**
5. **Founder-contributed IP.** Scout's social-layer substrate is the productionized continuation of KOAT.AI — a government-deployed manipulation detection tool founded by the Co-CEO. The IP and twenty years of virality-algorithm operating history transfer with the founder.

The surface architecture can be copied in 3–4 years for $10M+. By then, clients have compounded weight intelligence that does not transfer — and the falsifiability record cannot be backdated.

---

9. THE COMPETITIVE LANDSCAPE

The competitive landscape sits in seven categories. None of them produces what Vexorin produces. Bloomberg's surveillance architecture was built for a market that was single-surface; Vexorin complements Bloomberg by adding the multi-surface recognition layer rather than displacing the terminal.

Category	Representative vendors	Architecture	Surface	What Vexorin adds
Trade surveillance — exchange/regulator	NASDAQ SMARTS, Eventus (Validus), Trillium Surveyor, b-next	Rules-based	Order book / trade only	Cross-surface, agentic, on-prem
Bank-grade compliance surveillance	NICE Actimize, SteelEye	Rules + ML	Trade plus communications	Manifold geometry, manipulation purpose
Communications surveillance	Behavox, Smarsh	AI compliance	Voice / chat / email	Manipulation purpose (not comms), multi-surface
Fraud / anomaly ML	Featurespace, Quantexa	ML / entity resolution	Single-domain	Multi-surface, seven-family manipulation taxonomy
Crypto / on-chain	Solidus Labs, Chainalysis Markets, Elliptic, TRM Labs	Rules + ML	On-chain only	Cross-asset, full manifold
In-house builds at Tier-1 funds	Citadel, Millennium, Two Sigma, Renaissance, P72 internal teams	Custom, fund-specific	Variable	Pre-built and validated; lower TCO; transferable
Adjacent platforms (future)	Palantir Foundry; general-purpose agent platforms	General-purpose	Not domain-specific	Domain calibration, manipulation taxonomy

The white-space claim, sourced from literature. A literature search across the top ten published lite papers in agentic prediction analysis (TradingAgents, FinMem, FinVision, FinArena, AMA benchmark, AI-Trader, TimeSeriesScientist, DCATS, news-driven multi-agent papers, Autonomous Market Intelligence) and the top ten in manipulation detection and geometric representation (Hide-and-Shill, Sentiment Manipulation, P&D survey, Topological Complexity, Gidea-Katz TDA-in-finance, Topological Recognition crypto, Canadian Market Anomaly Detection, Null-Validated Topological Signatures, FX co-movements, ChainNet) returned no architecture fusing both. **Agentic synthesis plus geometric manifold representation has not been published.** The white space is in the literature itself, verifiable by any reader with access to arXiv and SSRN.

Vexorin was built for the market that exists now.

---

10. THE FUNDRAISE

10.1 The current round

$750K pre-seed at $3.5M pre-money, 17.65% dilution.

10.2 Use of funds

Team $300K · Engineering $220K · Data feed licensing $90K · Legal $60K · Ops $55K · GTM $25K. SR&ED Recovery $200K + Y1 Net Revenue $175K → Net Burn $375K. ~24 months runway post-close. Engineering allocation funds the parallel build streams: BRAIN aggregation engine, three additional agents (ECHO, ATLAS, CIPHER), and the Simulation Agent. Data feed licensing funds Scout’s 100-language coverage at customer scale.

10.3 The work this capital funds

Parallel execution across customer expansion AND data marketplace launch. Scout and Carlton deploy in 2–3 days into customer environments as production-ready agents (interaction model resembles deploying a language model). The capital funds: (1) the BRAIN aggregation engine and three additional agents going from architectural specification to running code, (2) the Simulation Agent that opens the Validated Data Marketplace as a Year 1 revenue stream, (3) onboarding the first cohort of data providers at $10K each + revenue share, (4) GTM execution across the three founder-owned distribution paths.

If we were generating signals at scale, we would not be at pre-seed. Live signal generation requires the ecosystem this round funds.

10.4 The six-month Pre-Seed-to-Seed timeline

The next six months execute against three goals: generate revenue, complete the ecosystem, build the cap table.

Today (Pre-Seed close). Scout and Carlton production-ready. LOIs in pipeline. CFO-underwritten projections. BRAIN architecture spec'd. 50-case library validated. Engineering team and Tier-1 agentic advisors in place. Founders Fund warm-intro path active.

Month 1–2. First three customers deploy Scout + Carlton on-prem (2–3 day deployments). Stream 1 ($50K/customer platform fees) and Stream 2 (agent consumption) begin flowing. Engineering parallel streams begin: BRAIN aggregation engine build, ECHO/ATLAS/CIPHER agent development, Simulation Agent design phase.

Month 3–4. First customer cohort fully operational. Cross-validation analytical pass complete (out-of-sample cosine numbers published in Prior Weight Matrix v1.1). Customer count expanding through warm-intro motion. Engineering progressing through parallel build streams.

Month 5. BRAIN aggregation engine operational. Three additional agents (ECHO/ATLAS/CIPHER) integrated into customer deployments. Simulation Agent in advanced build.

Month 6. Simulation Agent operational. Validated Data Marketplace opens. First data providers onboarded ($10K each + revenue share commencing). Stream 3 begins flowing. Full ecosystem deployed at customer sites.

Month 6+ to Series A close. Series A at $8–12M valuation. Cap table includes specialized fintech/regtech VCs, strategic angels, Tier-1 institutional finance LPs, plus the Series A lead bringing Tier-1 brand validation. The architecture has moved from "one auditable real-time detection (Iran)" to "deployed ecosystem with revenue across three streams."

10.5 The investor target shape

Primary path (highest conversion). Specialized fintech/regtech pre-seed VCs and angel-led rounds with strategic LPs in institutional finance. **Active warm-intro path.** Founders Fund via Engineering Advisor's direct connect (Trae Stephens is the obvious thesis fit — the Anduril-adjacent intelligence-community thesis maps directly to the team's KOAT.AI government deployment and British intelligence fieldwork). **Adjacent paths.** a16z American Dynamism (Boyle/Ulevitch); a16z fintech infrastructure (Rampell/Strange); Sequoia (selective).

10.6 Cap table construction

This pre-seed builds the foundation cap table. Series A brings the Tier-1 lead and the brand validation that compounds through Series B. The pre-seed cohort is not a placeholder for "real" investors who arrive later — it is the foundation layer of a deliberately constructed cap table where each round adds capability the company needs at that stage.

10.7 The three-deck strategy

Deck 1 — Outreach (14 slides, cold-outreach optimized, warm-intro cover variant). Earns the meeting. **Deck 2 — First Call** (~25–35 slides, technical depth). Earns the diligence. **Deck 3 — Due Diligence** (full technical and commercial depth, data room companion). Earns the term sheet.

---

11. KEY MESSAGES FOR ANY EXTERNAL COMMUNICATION

When generating any deck, pitch, email, or external communication, the following claims and phrasings are the canonical version. Other phrasings should defer to these.

- *Recognition is the new alpha.* (The single canonical tagline.)
- *They see the data. We see the patterns.* (The canonical sub-tagline.)
- *We are building the recognition substrate for institutional capital. Not a feature. Not a vendor. The architecture the next decade is built on.* (The platform-thesis frame.)
- *The cascade is the shape.* (The architectural primitive.)
- *The base image is real before the cascade can extend.* (The anti-hallucination defense.)
- *Recognition is the bottleneck. Prosecution is downstream of recognition.* (The market-thesis frame.)
- *The slow accumulation of confirmed events is not a methodology limit — it is the category's structural rarity, the same rarity Vexorin exists to address.* (The CALIBRATE-rarity defense.)
- *Production-ready. Pre-revenue by design. Demand validated by the LOIs in the pipeline.* (The pre-seed honesty frame.)
- *If we were generating signals at scale, we would not be at pre-seed.* (The clinching pre-seed line.)
- *Manipulation hides in fragments. Together, they form a pattern. We see the pattern. Vexorin.* (The closing manifest.)

---

12. WHAT THIS DOCUMENT IS NOT

This document is not the BRAIN Litepaper. Architectural depth, conviction-aggregation math, signal contracts, family-specific weight overrides, the Iran 2026 cascade trace, and the deployment / governance technical detail all live there. This document complements rather than duplicates.

This document is not the Prior Weight Matrix v1.0. Per-agent weight derivations, family-centroid cosine similarity tables, and the falsifiability test methodology live there.

This document is not the 50-Case Manifold Pattern Analysis. The case-by-case retrospective reconstruction lives there.

This document is not the Citations Audit v2.0. Per-citation verification status and audit trail live there.

This document is not a public-facing brochure. It is internal canonical reference material, calibrated for advisors, prospective investors at the diligence stage, and team members generating external communications. The deck and the public website carry compressed and audience-tuned versions of the content here.
